

Aalto University
School of Science
Master's Programme in Computer, Communication and Information Sciences

Artem YUSHKOVSKIY

Automated Analysis of Weak Memory Models

Master's Thesis
Espoo, ???2018

Supervisor: Assoc. Prof. Keijo Heljanko
Instructor:

Aalto University
 School of Science

 Master's Programme in Computer, Communication and In-
 formation Sciences

 ABSTRACT OF
 MASTER'S THESIS

Author:	Artem YUSHKOVSKIY		
Title:	Automated Analysis of Weak Memory Models		
Date:	?.?.2018	Pages:	vi + 27
Professorship:		Code:	AS-116
Supervisor:	Assoc. Prof. Keijo Heljanko		
Instructor:			
<p>In id fringilla velit. Maecenas sed ante sit amet nisi iaculis bibendum sed vel elit. Quisque eleifend lacus nec ipsum lobortis ornare. Nam lectus diam, facilisis eget porttitor ac, fringilla quis massa. Phasellus ac dolor sem, eget varius lacus. Sed sit amet ipsum eget arcu tristique aliquam. Integer aliquam velit sit amet odio tempus commodo. Quisque commodo lacus in leo sagittis vel dignissim quam vestibulum. Cras fringilla velit et diam dictum faucibus. Pellentesque at eros non mauris auctor euismod. Nullam convallis arcu vel lectus sollicitudin rutrum. Praesent consequat, nisl at pretium posuere, neque arcu dapibus lacus, ut sollicitudin elit velit ultricies libero. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae;</p> <p>Nulla semper hendrerit molestie. Pellentesque blandit velit sit amet est vestibulum faucibus. Nullam massa turpis, venenatis non mollis fringilla, mattis et diam. Fusce molestie convallis elementum. Morbi nec lacus dapibus arcu mollis gravida. Aliquam erat volutpat. Nam vitae magna nunc. Nunc ut ipsum at massa porttitor vestibulum. Praesent diam lorem, ultrices nec vestibulum id, volutpat nec lacus.</p>			
Keywords:	Thesis template, master’s thesis		
Language:	English		

Aalto-yliopisto
 Perustieteiden korkeakoulu
 ???

???

Tekijä:	Artem YUSHKOVSKIY		
Työn nimi:	?		
Päiväys:	???.2018	Sivumäärä:	vi + 27
Professuuri:	?	Koodi:	AS-116
Valvoja:	Assoc. Prof. Keijo Heljanko		
Ohjaaja:	<p>Cras tincidunt bibendum erat, vel tincidunt diam porttitor aliquam. Donec sit amet urna non felis placerat pharetra. Aenean ultrices facilisis nulla vitae semper. Nullam non libero quis dui fermentum aliquam id vel eros. Praesent elementum tortor quis sem congue iaculis sit amet eget nisl. Quisque erat tortor, condimentum eu volutpat et, blandit et augue. Phasellus erat turpis, pretium non feugiat id, posuere id velit. Vestibulum ut sapien felis, quis convallis dui.</p> <p>In elementum est eu nulla hendrerit feugiat. In sodales diam vel lacus cursus tincidunt. Morbi nibh dui, imperdiet non vestibulum non, dignissim id risus. Sed sollicitudin neque lectus, porttitor sollicitudin elit. Nulla facilisi. Nullam in ante eu mi suscipit sollicitudin. Sed est velit, gravida facilisis varius eget, tempus sed urna. Aliquam erat volutpat. Nam semper condimentum nisi. Nullam scelerisque, metus nec sodales vulputate, purus augue venenatis urna, sit amet mattis turpis nisl ac metus. Mauris nec odio ut neque condimentum vulputate vel in turpis. Nulla facilisi. Nulla id tellus sapien, vitae blandit lorem.</p>		
Asiasanat:	Diplomityöpohja		
Kieli:	Englanti		

Acknowledgements

In id fringilla velit. Maecenas sed ante sit amet nisi iaculis bibendum sed vel elit. Quisque eleifend lacus nec ipsum lobortis ornare. Nam lectus diam, facilisis eget porttitor ac, fringilla quis massa. Phasellus ac dolor sem, eget varius lacus. Sed sit amet ipsum eget arcu tristique aliquam. Integer aliquam velit sit amet odio tempus commodo. Quisque commodo lacus in leo sagittis vel dignissim quam vestibulum. Cras fringilla velit et diam dictum faucibus. Pellentesque at eros non mauris auctor euismod. Nullam convallis arcu vel lectus sollicitudin rutrum. Praesent consequat, nisl at pretium posuere, neque arcu dapibus lacus, ut sollicitudin elit velit ultricies libero. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae;

Espoo, ???.2018

Artem YUSHKOVSKIY

Abbreviations

LI	Lorem Ipsum
ABC	Quisque et mi lacus, nec porta ante.
DEF	Proin pellentesque accumsan laoreet

Contents

Abbreviations	v
1 Introduction	1
1.1 Thesis structure	4
2 Memory Model-Awared Program Analysis	5
2.1 The event-based program representation	5
2.1.1 Events	6
2.1.2 Relations	7
2.1.3 Executions	7
2.1.4 Weak Memory Models	9
2.2 The CAT language	9
2.3 Some examples of WMM	9
3 The Portability of Concurrent Software	10
3.1 The model checking problem	10
3.2 The portability as a bounded reachability problem	11
3.2.1 The control-flow encoding	12
3.2.2 The data-flow encoding	14
3.2.3 The memory model encoding	15
4 The input language	17
5 The mousquetaires: implementation	19
5.1 Requirements	19
5.2 Program Components	20
5.3 Parsing the input language as YTree	20
5.4 YTree to XGraph event converter	21
5.4.1 Loop unrolling	21
5.5 XGraph to ZFormula (SMT) encoder	22

5.6	Optimisations	22
6	Evaluation	23
6.1	Comparison with PORTHOS	23
6.1.1	Unique Features	23
6.1.2	Performance	23
6.2	Comparison with HERD	23
6.2.1	Unique Features	23
6.2.2	Performance	23
7	Summary	24
	Bibliography	25
	Appendices	25
A	The ANTLR grammar of input language used in PORTHOS	26

Chapter 1

Introduction

Most modern computer systems contain large parts that operate concurrently. Though parallelisation of the system can improve its performance drastically, it opens numerous of problems connected to correctness, robustness and reliability, which makes the concurrent program design one of the most difficult problems of programming [McK17].

Traditionally, studies related to concurrent programming concern on more fundamental theoretical questions of designing race-free and lock-free parallel algorithms, asynchronous data structures and synchronisation primitives of a programming language. Unfortunately, when it comes to the real-world concurrent programs, the algorithmic level of abstraction is not enough for guaranteeing their properties of correctness and reliability. The reasons of this fact lie in the code optimisations that both compiler and hardware perform in order to increase performance as much as possible. For instance, Figure 1.1 provides simple example of reachability of the state $(0:EAX=0 \wedge 1:EAX=0)$ on x86 machines (such little examples that illustrate specific behaviour of a WMM are called *litmus tests*). This state is allowed because in x86 architecture each processor may cache the write to shared memory variable into its local write buffer, so that they do not become visible by other processes immediately. In the example, the write `'MOV [x], 1'` performed by process P0 stores value 1 to the shared variable [x] into the write buffer of process P0. Meanwhile, the write cache of the process P1 may not have updated version of the variable [x], neither may have the main memory, so that the read `'MOV EBX, [x]'` performed in the process P1 may read the initial value 0 even if this variable has been already updated in another thread. These problems have lead to the need for formalisation of

{ x=0; y=0; }	
P0	P1
MOV [x],1	MOV [y],1
MOV EAX,[y]	MOV EAX,[x]
exists (0:EAX=0 /\ 1:EAX=0)	
x86-TSO: allow	

Figure 1.1: Store buffering (SB): a litmus test on write-read reordering allowed under the x86-TSO weak memory model

semantics of memory operations within different concurrent architectures defined by *weak memory models* (WMM).

Research of weak memory models firstly aims to *formalise* develop the formal approach of understanding programs with respect to weak memory models which is systematic, sound and complete. The first (and so far the only) such a framework was presented in 2010 [Alg10]. Secondly, researchers work on extracting the WMMs for hardware architectures from existing implementations of from their specifications, that are written in natural language and thus suffer from ambiguities and incompletenesses. Over last decade the memory models have been extracted for most mainstream multiprocessor architectures, such as x86-TSO (*Total Store Order*) model for x86 architecture formalised in 2009 [OSS09], much more relaxed memory model for Power and ARM architectures [SSA⁺11] [AFI⁺09], for Alpha ? and Sparc ? memory models were defined in the specification. Moreover, most modern high-level programming languages rely on relaxed memory model. Thus, the memory model for Java that is based on the *happens-before* principle [Lam78] was introduced in J2SE 5.0 in 2004 [MPA05]; the C++11 standard ? has introduced the set of hardware-independent synchronisation fences and atomic operations, whenever the C++17 memory model is based on the relation *strongly happens-before*. Thirdly, important research direction targets the problem of verifying (or at least finding bugs in) existing software systems with respect to WMMs. Perhaps, the most notable work in this field is the project for defining the Linux kernel memory model, which is being actively developing these days [MAM⁺17].

The first memory model for concurrent systems was formulated by Leslie Lamport back in 1979 [Lam79]. This memory model, called the *sequential consistency* (SC), allows only those executions (interleavings) that produce the same result as if the operations had been executed by single process. This means that the order of operations executed by a

process is strictly defined by the program it executes. The SC model does requires the write to a shared variable performed in one process to become visible by all other processes not instantly, but simultaneously. This means each process communicates to the shared memory directly, without local buffering. Another important requirement of SC memory model is that it forbids memory operations reordering within single process (the order is strictly defined by the program).

The SC model is considered to be the strong memory model in the sence that it provides strong guarantees regarding the ordering and caused effect of memory operations. Different relaxations of this model lead to the class of *weak memory models* (WMM). They specify how threads interact through shared memory, when a write becomes visible to other threads and what value a read can return. Therefore, WMMs serve as set of guarantees made by designers of execution environment (hardware, programming language, compiler, database, operation system, etc.) to programmers on which behaviours of their concurrent code they may expect.

Although weak memory studies is rather young research area, there exist frameworks and tools for exploring WMMs and examining simple programs with respect to the them. The state-of-the-art tool is diy (for *do it yourself*), developed by the researchers from INRIA institute, France and University of Cambridge, UK. The diy¹ is a software suite for designing and testing weak memory models. It is firstly released back in 2010, and since that time it remained to be the only tool for testing weak memory models. The diy consists of several modules: the litmus tests generators diy, the litmus tests concrete executor litmus that runs tests on a physical machine while collecting its behaviours, and the weak memory models simulator herd that implements reachability analysis for exploring states reachable under specified WMM.

All the diy tools work only with single memory model, however, in real life we face serious engineering problems involving necessity to model more than one execution environment. One of these problems is the *portability* of the program from one hardware architecture to another. A program written in a high-level language is then compiled for different hardware. Even if all the compiler optimisations were disabled (which is rare case nowadays), the behaviour of two compiled versions of the same program may differ due to differences between hardware memory models. As the result, a program compiled under the platforms T can reach states that are unreachable on

¹Project web page: diy.inria.fr

the platform S , which is a *portability bug* from the source platform S to the target platform T [LFH⁺17].

The first tool that performs the WMM-awared portability analysis is PORTHOS introduced in April 2017 [LFH⁺17]. This tool reduces described problem to a bounded reachability problem, which can be solved with help of an SMT-solver. This approach allows to capture symbolically the semantics of analysing program and both weak memory models into single SMT-formula, augmented by the reachability assertion. As most modern SMT-solvers are efficient enough to be able to operate the state space of size millions of variables bounded by millions of constraints ([MZ09]), the used method can be applicable in solving the real-world problems.

Current work aims to rework the proof-of-concept tool PORTHOS by extending the input language, which currently represents the minimum subset of C, and revising the general architecture of the tool in order to enhance performance, reliability and mantainability. We called the new tool **the mousquetaires framework** as now it not only performs the portability analysis, but can serve as the basis for SMT-solver driven weak memory model analysis.

1.1 Thesis structure

The Chapter 2 gives more detailed description of the weak memory model-awared analysis and provides description of memory models for some common architectures (x86, ARM and POWER, Sparc, ???). Chapter ...

Chapter 2

Memory Model-Aware Program Analysis

In general, analysis of concurrent programs with respect to axiomatic memory models is performed in several stages. Firstly, the control-flow and data-flow of a program is encoded as the set of possible *candidate executions*. Obtained model of the program is called an anarchic semantics, which is a truly parallel semantics with no global time that describes all possible computations with all possible communications [ACM16]. Thereafter, the anarchic semantics is constrained by the *weak memory model* specification which is a set of axiomatic constraints for filtering out executions inconsistent in particular architecture.

2.1 The event-based program representation

The classical approach to model concurrent programs is to use the *global time*, a single order of interleavings among all events happened in different threads. However these models are easy to understand, it may be impossible to process *all* possible states, number of which is exponentially large. However, there exist some equivalence classes such that the result of execution different interleavings from single equivalence class is the same (for instance, computations performed by a processor locally do not affect the global state). One such model is the *event-based* representation of a program, which models a program as a directed graph (the *event-flow graph*). The vertices of such a graph represent *events* (simple low-level instructions; see Section 2.1.1), and edges represent *relations* over the events

(see Section 2.1.2). Thus, a candidate execution constitutes the set of guesses which event (from the set of possible next events) is to be executed.

2.1.1 Events

A *memory event* $e_m \in \mathbb{E}$ represents the fact of access to the memory. Since memory is the crucial low-level resource shared by multiple processes, most relations are defined over memory events. The processes can access a shared memory location (denoted by l_i , for *location*), or a local one (denoted by r_i , for *register*). A memory event can access at most one shared memory location, high-level instructions that address more than one shared variable must be transformed into a sequence of events. A memory event is specified by its direction with respect to the shared variable, its location $\text{loc}(e_m)$, its processor label $\text{proc}(e_m)$, and a unique event label $\text{id}(e_m)$ [Alg10].

The set of memory events \mathbb{M} is divided into write events \mathbb{W} (that write values to shared-memory locations) and read events \mathbb{R} (that read values stored in shared-memory locations). We add a restriction that each memory event uses at most one shared location, so that the write instruction $i = \text{write}(l_1, l_2)$, that encodes the write from the shared location l_2 to the shared location l_1 , is represented as two consequent events $e_1 = \text{load}(r_1 \leftarrow l_2)$; $e_2 = \text{store}(l_1 \leftarrow r_1)$. Also, it is important to separate the set of initial write events $\mathbb{IW} \subset \mathbb{W}$ that perform initialisation of program variables.

A *computation event* $e_c \in \mathbb{C} \subseteq \mathbb{E}$, represents a low-level assembly computation operation performed solely on local-memory arguments. An example of computation event may be the event $e_c = r_1 \leftarrow \text{add}(r_2, 1)$ that writes the sum of values stored in register r_2 and constant 1 (which is modelled as a register as well) to the register r_1 . For modelling branching statements, we distinguish the set $\mathbb{C}_1 \subseteq \mathbb{C}$ of *predicative* computation events (also called as *branching events*), that are evaluated as a boolean value.

The synchronisation instructions (fences) cause the *barrier events*, that do not perform any computation or memory value transfer, instead, they add new relations to the program model that restrict the set of allowed behaviours. Semantically, a fence may serve as a synchronisation barrier, or as a flush of local memory caches, etc.

2.1.2 Relations

The relation $\xrightarrow{r}: \mathbb{E} \times \mathbb{E}$ is a binary function over events (set of pairs of events). There are two kinds of relations between events: *basic relations* that capture semantics of the program, and *derived relations* that are defined from the basic relations and events in the weak memory model specification. Constraints over relations that are specified by weak memory models are defined as requirements of acyclicity, irreflexivity or emptiness of specific relations [ACM16].

- The *control-flow* of a program is defined by the *program-order* relation $po \subset \mathbb{E} \times \mathbb{E}$, which represents the total order of events of same process. For instance, if the instruction i_1 generates the event e_1 and the instruction i_2 follows i_1 and generates the event e_2 , then $e_1 \xrightarrow{po} e_2$.
- The *data-flow* of a program is defined by *communication relations*: the *read-from* relation $rf \subset \mathbb{W} \times \mathbb{R}$ that maps each write event to the read event that reads its value, the *coherence order*-relation $co \subset \mathbb{W} \times \mathbb{W}$ defines the total order on writes to the same location across all processes.
- Events from the same process are related by the *scope relation* $sr \subset \mathbb{E} \times \mathbb{E}$.

The derived relations can be produced using the following operations over relations [ACM16]:

- *the union* of two relations $r1$ and $r2$ is $r1 \mid r2$
- *the intersection* of two relations $r1$ and $r2$ is $r1 \& r2$
- *the difference* of two relations $r1$ and $r2$ is $r1 \setminus r2$
- *the sequence*¹ of two relations $r1$ and $r2$ is $r1 ; r2$

2.1.3 Executions

The semantics of a concurrent program is represented by the set of allowed executions. A *candidate execution* (trace, run) of a program is an ordered set

¹The sequence of two relations $r1$ and $r2$ is defined as the set of pairs (x, y) such that there exists an intervening z , such that $(x, z) \in r1$ and $(z, y) \in r2$

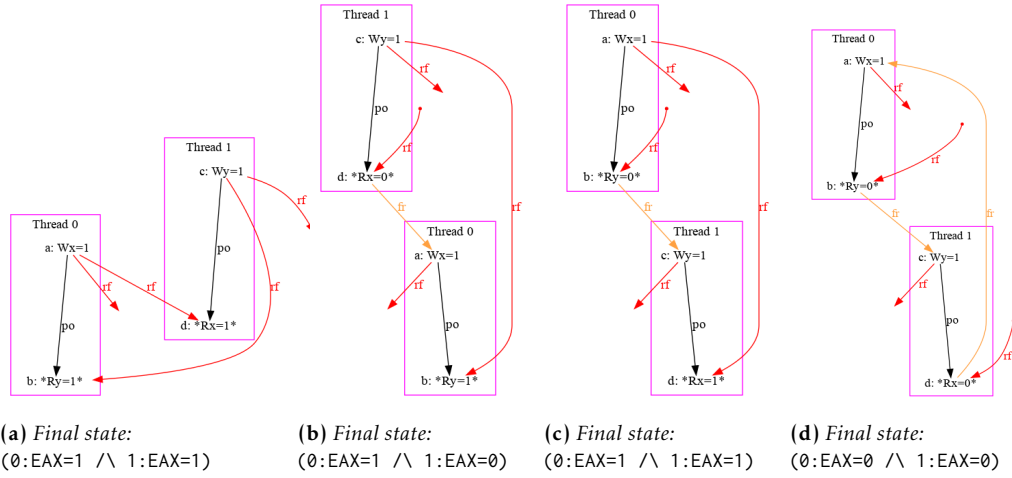


Figure 2.1: Candidate executions for the litmus test Example 1.1

of events defined by po - and rf -relations and set of final writes to a given memory location [AMT14].

An execution is uniquely defined by the set \mathbb{X} of events have been executed in each thread (the *control-flow* of a program), and the relations rf and co [Alg10]. As it was shown in [WBS⁺17], it is enough for memory models to constrain the executions independently instead of constraining the program as a whole. Figure 2.1 illustrates four possible candidate executions for the litmus-test Example 1.1 (the pictures are generated by the herd7 tool, version 7.47). Since there are no conditional jumps, the po-relation is defined and we do not need to guess it. Since each thread performs single write followed by a single read, the co-relation is also defined (it relates the initial write event with the write event to the same location). Thus, there are only four possible executions defined by the choice of rf -relation. The candidate executions pictured in Figures 2.1a–2.1c are consistent both under strong memory model SC and under relaxed memory models x86-TSO, Power, ARM, and others. However, the execution shown in Figure 2.1c is still consistent under relaxed-memory architectures, but it becomes inconsistent under SC architecture as it forbids cycles over $fr \cup po$.

2.1.4 Weak Memory Models

2.2 The CAT language

One common way to define a weak memory model in a systematic way is to formalise it in *the CAT language* proposed in [ACM16]. The CAT language allows to axiomatically define new relations, new fences and restrictions over relations.

- briefly some hw memory models: X86-TSO, Alpha, POWER, ;
- language memory models: Java, C++;
- library-level kernel memory model, ref to github with tests

2.3 Some examples of WMM

```
//axioms of TSO wmm
// example of sets: rf, co, ... for the code snipped used before
```


Chapter 3

The Portability of Concurrent Software

As it has been discussed in Chapter 1, the program may behave differently when compiled for different parallel hardware architectures. This may cause the portability bugs, the behaviour that is allowed under one architecture and forbidden under another. In this Chapter, we describe the general task of analysing the concurrent software portability as a *bounded reachability* problem, which in turn can be reduced to the satisfiability modulo theories (SMT) problem [LFH⁺17].

3.1 The model checking problem

The classical model checking algorithms explore the state space of an abstract automata or a transition system in order to find states that violate the specification. The general schema of model checking is the following: firstly, the analysing system is being represented as a transition system, a finite directed graph with labeled nodes representing states of the system such that each state corresponds to the unique subset of atomic propositions, that characterise the behavioral properties of each state. Then, the system constraints are being defined in terms of a modal temporal logic with respect to the atomic propositions. Commonly, the Linear Temporal Logic (LTL) or Computational Tree Logic (CTL), along with their extensions, are used as a specification language due to the expressiveness and verifiability of their statements. In the described schema, the model checking problem is reducible to the reachability analysis, an iterative process of a systematic

exhaustive search in the state space. This approach is called *unbounded model checking (UMC)*.

However, all model checking techniques are exposed to the *state explosion problem* as the size of the state space grows exponentially with respect to the number of state variables of the system. In case of modeling concurrent systems, this problem becomes much more considerable due to exponential number of possible interleavings of states. Therefore, the research in model checking over past 40 years was aimed at tackling the state explosion problem, mostly by optimising search space, search strategy or basic data structures of existing algorithms.

One of the first technique that optimises the search space considerably major was the symbolic model checking with binary decision diagrams (BDDs). Instead of by processing each state individually, in this approach the set of states is represented by the BDD, efficient data structure for performing operations on large boolean formulas [CKN⁺12]. The BDD representation can be linear of size of variables it encodes if the ordering of variables is optimal, otherwise the size of BDD is exponential. The problem of finding such an optimal ordering is known as NP-complete problem, which makes this approach inapplicable in some cases.

The other idea is to use satisfiability solvers for symbolic exploration of state space [CBR⁺01]. In this approach, the state space exploration consists of sequence of queries to the SAT-solver, represented as boolean formulas that encode the constraints of the model and the finite path to a state in the corresponding transition system. Due to the SAT-solver. This technique is called *bounded model checking (BMC)*, because the search process is being repeated up to user-defined bound k , which may result to incomplete analysis in general case. However, there exist numerous techniques for making BMC complete for finite-state systems (e.g., [Sht00]).

3.2 The portability as a bounded reachability problem

In general, the BMC problem aims to examine the reachability of the "undesirable" states of a finite-state system. Let $\vec{x} = (x_1, x_2, \dots, x_n)$ be a vector of n variables that uniquely distinguishes states of the system; let $Init(\vec{x})$ be an *initial-state predicate* that defines the set of initial states of the system; let $Trans(\vec{x}, \vec{x}')$ be a *transition predicate* that signifies whether there the transition from state \vec{x} to state \vec{x}' is valid; let $Bad(\vec{x})$ be a *bad-state predicate*

that defines the set of undesirable states. Then, the BMC problem, stated as the reachability of the undesirable state withing k steps is formulated as following: $\text{SAT}(\text{Init}(\vec{x}_0) \wedge \text{Trans}(\vec{x}_0, \vec{x}_1) \wedge \cdots \wedge \text{Trans}(\vec{x}_{k-1}, \vec{x}_k) \wedge \text{Bad}(\vec{x}_k))$.

The portability program may also be formulated as a reachability problem, where the undesirable state is the state reachable under the target \mathcal{M}_T memory model and unreachable under the source memory model \mathcal{M}_S . However, unlikely the BMC problem, the portability analysis does not require to call the SMT-solver repeatedly, since (imperative) programs may be converted as acyclic state graph (by reducing the loops, see Section 5.4.1) and the *Trans* predicate may be stated only for the final state of a program.

As we have said before, a program P is called portable from the source architecture (defined as a memory model) \mathcal{M}_S to the target architecture \mathcal{M}_T if *all* executions consistent under \mathcal{M}_T are consistent under \mathcal{M}_S [LFH⁺17]:

Definition 3.2.1 (Portability). Let $\mathcal{M}_S, \mathcal{M}_T$ be two weak memory models. A program P is portable from \mathcal{M}_S to \mathcal{M}_T if $\text{cons}_{\mathcal{M}_T}(P) \subseteq \text{cons}_{\mathcal{M}_S}(P)$

Note, that the formulation of portability requirements against *executions* is strong enough, as it implies the portability against *states* (the *state-portability*) [LFH⁺17].

The result SMT formula ϕ of the portability problem should contain encodings of control-flow (ϕ_{CF}) and data-flow (ϕ_{DF}) of the program, and assertions of both memory models: $\phi = \phi_{CF} \wedge \phi_{DF} \wedge \phi_{\mathcal{M}_T} \wedge \phi_{\neg \mathcal{M}_S}$. If the formula is satisfiable, there exist a portability bug.

3.2.1 The control-flow encoding

The control-flow of a program is represented in the *control-flow graph*, a directed acyclic connected graph with single source and multiple sink nodes, obtained by the *loop unrolling* (see Section 5.4.1). In control-flow graph, there are two types of edges: *primary edges* that denote unconditional jumps or if-true-transitions (pictured with solid lines), and *alternative edges* that denote if-false-transitions (pictured with dotted lines). Each node on graph can have either one successor (primary) or two successors (primary and alternative); only computation events can serve as a branching point). However, each merge node can have any positive number of predecessors, where each edge may be either primary or alternative.

While working on the mousquetaires, we applied some modifications of the encoding scheme for the control-flow. The changes are conditioned by the need to be able to process an arbitrary control-flow produced by

conditional and unconditional jumps of C language. For that, we compile the recursive abstract syntax tree (AST) of the parsed C-code to the plain (non-recursive) event-flow graph. We show that the new encoding is smaller than the old one used in PORTHOS since it does not produce new variables for each high-level statement of the input language. For instance, in old scheme the control-flow of the sequential instruction $i_1 := i_2; i_3$ was encoded as $\phi_{CF}(i_2; i_3) = (cf_{i_1} \Leftrightarrow (cf_{i_2} \wedge cf_{i_3})) \wedge \phi_{CF}(i_2) \wedge \phi_{CF}(i_3)$, and control-flow of the branching instruction $i_1 := \{c?i_2 : i_3\}$ was encoded as $\phi_{CF}(c?i_2 : i_3) = (cf_{i_1} \Leftrightarrow (cf_{i_2} \vee cf_{i_3})) \wedge \phi_{CF}(i_2) \wedge \phi_{CF}(i_3)$ (here we used the notation of C-like ternary operator ' $x?y:z$ ' for defining the conditional expression 'if x then y else z '). In contrast, the new scheme implemented in *mousquetaires* firstly compiles the recursive high-level code into the linear low-level event-based representation, that is then encoded into an SMT-formula. The encoding of branching nodes depends on the *guards*, the value of conditional variable on the branching state, which in turn is encoded as data-flow constraint (see further in current Chapter).

Let $\mathbf{x} : \mathbb{E} \rightarrow \{0, 1\}$ be the predicate that signifies the fact that the event has been executed (and, consequently, has changed the state of the system). Let $\mathbf{v} : \mathbb{C} \rightarrow \mathbb{N}$ be the function that returns the value of the computation event that will be computed once the event is executed (strictly speaking, it returns the *set* of values determined by the *rf*-relation; see Chapter ? for the relations encoding). We distinguish the function $\mathbf{v}_p : \mathbb{C}_1 \rightarrow \{0, 1\}$ that evaluates the predicative computation event. In the result formula, all symbols $\mathbf{x}(e_i)$ and $\mathbf{v}(e_i)$ are encoded as boolean variables.

Consider the following possible mutual arrangement of nodes in a control-flow graph:




Figure 3.1: Linear and non-linear cases of control-flow graph

For listed cases, below we propose the encoding scheme that uniquely encodes each node of graph and allows to encode partially executed program. The Equation 3.1 for encoding the sequential control-flow represented in Figure 3.1a reflects the fact that the event e_2 can be executed iff the event

e_1 was executed. The Equation 3.2 for encoding the branching control-flow depicted in Figure 3.1b allows only following executions: $\{\emptyset, (e_1), (e_1 \rightarrow e_2), (e_1 \rightarrow e_3)\}$. In encoding 3.3 of the merge-point represented in Figure 3.1c, the event e_k is executed if either of its predecessors was executed, regardless of type of the transition.

$$\phi_{CF_{seq}} = \mathbf{x}(e_2) \rightarrow \mathbf{x}(e_1) \quad (3.1)$$

$$\begin{aligned} \phi_{CF_{br}} = & [\mathbf{x}(e_2) \rightarrow \mathbf{x}(e_1)] \wedge [\mathbf{x}(e_3) \rightarrow \mathbf{x}(e_1)] \wedge \\ & [\mathbf{x}(e_2) \rightarrow \mathbf{v}(e_1)] \wedge [\mathbf{x}(e_3) \rightarrow \neg \mathbf{v}(e_1)] \wedge \\ & \neg[\mathbf{x}(e_2) \wedge \mathbf{x}(e_3)] \end{aligned} \quad (3.2)$$

$$\phi_{CF_{merge}} = \mathbf{x}(e_k) \rightarrow \left(\bigvee_{e_p \in \text{pred}(e_k)} \mathbf{x}(e_p) \right) \quad (3.3)$$

3.2.2 The data-flow encoding

While encoding the data-flow constraints, we use the *Static Single-Assignment (SSA) form* in order to be able to capture an arbitrary data-flow into a single SMT-formula. The SSA form requires each variable to be assigned only once within entire program. That is, on each state of the program a variable (either local or shared) is mapped to its indexed reference, which represents a new boolean variable in the formula. On merge-points, the ϕ -function is emulated as the disjunction of values of the variable that is computed as the *reaching definition*. For that, we carry the map that for each event stores mapping from all declared variables to their SSA-indices. All indices are globally unique within single program analysis.

As it is implemented in [LFH⁺17], the SSA-indices of variables are computed in accordance with the following rules: (1) any access to a shared variable (both read and write) increments its SSA-index; (2) only writes to a local variable increment its SSA-index (reads preserve indices); (3) no access to a constant variable or computed expression changes their SSA-index. These rules determine the following encoding of load, store and computation events within single thread:

$$\phi_{DF_{e=\text{load}(r \leftarrow l)}} = \mathbf{x}(e) \rightarrow (r_{i+1} = l_{i+1}) \quad (3.4)$$

$$\phi_{DF_{e=\text{store}(l \leftarrow r)}} = \mathbf{x}(e) \rightarrow (l_{i+1} = r_i) \quad (3.5)$$

$$\phi_{DF_{e=\text{compute}(\dots)}} = \mathbf{x}(e) \rightarrow \mathbf{v}(e) \quad (3.6)$$

Unlike the PORTHOS tool, the *mousquetaires* reuses the algorithm of converting the program into the SSA form was also adapted for processing the plain flow-graph data structure. That is, for each node we keep track the set of its predecessors, and then, while traversing the graph in topological order, we propagate the SSA-indices for each variable defined so far, and update these indices according the rules listed above.

The time of described algorithm is linear of the size of event-flow graph since it performs only single traverse of the graph. Notwithstanding the overhead of storing (or, equivalently but undesirably, computing with linear time) the predecessors of each node in order to be able to convert the program into an SSA form, the time of such a transformation reduced comparing to the algorithm implemented in PORTHOS, which recomputed SSA-indices recursively for each instruction.

Moreover, the data-flow encoding scheme used in PORTHOS faced the problem when the branching statements had different number of assignments of certain variable in their branches. //TODO: why? why don't we have this problem now?

As it has been stated before, the *rf*-relation links the parts of data-flow stored in equivalence assertions over the SSA-variables. The encoding of this linkage left untouched as it is implemented in PORTHOS: for each pair of events e_1 and e_2 linked by the *rf*-relation, we add the following constraint:

$$\phi_{DF_{mem}}(e_1, e_2) = rf(e_1, e_2) \rightarrow (l_i = l_j) \quad (3.7)$$

where the variable of location l is mapped to the SSA-variable l_i for event e_1 , and to the SSA-variable l_j for event e_2 ; and the predicate $rf(e_1, e_2)$ is encoded as a boolean variable, which itself equals *true* if e_2 reads the shared variable that was written in e_1 .

3.2.3 The memory model encoding

The encoding of a weak memory model itself does not depend on the analysing program. The problem here...

recursive definitions and Kleene iteration

As it is described in [LFH⁺17], the is based ...

Chapter 4

The input language

The PORTHOS tool used a small subset of C as an input language as described in Figure 4.1 [LFH⁺17]:

```
<prog> : program <thrd>*  
      ;  
<thrd> : thread <tid> <inst>  
      ;  
<inst> : <atom>  
      | <inst> ; <inst>  
      | while <pred> <inst>  
      | if <pred> then <inst> else <inst>  
      ;  
<atom> : <reg>  $\leftarrow$  <exp>  
      | <reg>  $\leftarrow$  <loc>  
      | <loc> := <reg>  
      | mfence  
      | sync  
      | lwsync  
      | isync  
      ;
```

Figure 4.1: *Syntax of an input language of PORTHOS*

The input language of PORTHOS supported recursive definitions of control-flow instructions, that were directly encoded into SMT-formula. In *mousquetaires*, the set of supported instructions is extended in order to support expressions valid in C. Also, for the purpose of simplicity, all kinds of fence instructions and different types of variable assignments (shared variable assignment ‘:=’ or local variable assignment ‘ \leftarrow ’) were defined directly in the grammar for PORTHOS, whereas the semantics of function invocations need to be determined dynamically during the interpretation (see Chapter ??).

Both PORTHOS and *mousquetaires* use the parser generator ANTLR [Par13], the powerfull language processing tool. The (hand-written) ANTLR grammar used in PORTHOS is available at Appendix A, whereas the *mousquetaires* uses the grammar of C language proposed in the C11 standard [ISO11] (the ANTLR grammar can be found in the official ANTLR repository on GitHub ¹).

¹<https://github.com/antlr/grammars-v4>

Chapter 5

The mousquetaires: implementation

This Chapter describes the architecture of the tool mousquetaires.

The programming language choice for mousquetaires was made in favour of java in order to reuse some parts of its predecessor PORTHOS that was written in java. However this language does not show best results in performance benchmarks (comparing to, for example, C++) **TODO**, the performance cornerstone of mousquetaires (as well as any other SMT-based code analyser) is the phase of solving the SMT-formula, which is left to the third-party SMT-solver called from mousquetaires via java API. However, considering the perspective of using mousquetaires as a static analyser for real-world programs, we also must take into account memory problems at both encoding and solving stages.

5.1 Requirements

The main reason for starting the work on re-implementing the PORTHOS tool was the need to extend input language to be able to process real C programs in perspective. The existing PORTHOS architecture mixes high-level recursive *instructions* (statements of C) with low-level non-recursive *events*, to which the high-level instructions are compiled (see classes of package ‘dartagnan.program’ in the PORTHOS project ¹). Both instructions and events are then encoded into the SMT-formula. However, this approach is poorly extensible as to adding support for a high-level control-flow instruction (say, the do-while loop in C) requires making changes in many places of code from parser to encoder. Moreover, we encounter serious problems

¹Project web page: <http://github.com/hernanponcedeleon/PORTHOS>

when we need to add support for control-flow jumps (as `continue`, `break`, `goto` in C).

Therefore, while planning the architecture of *mousquetaires*, we decided to clearly separate high-level recursive code representation as Abstract Syntax Tree (AST) and low-level compiled event-based representation as a flow-graph. Such a modular architecture allows to use multiple input-language parsers and convert parsed syntax trees to our AST, thus having support for multiple languages (for instance, the original input language of PORTHOS tool, two syntaxes of litmus tests used by *herd*, and even assembly language for any supported architecture).

The target **requirements** for the new tool was set up as following (in descending order of priority):

1. stability and transparency (KISS principle)
 - choice in favour of simplicity and readability at every step;
 - using software design patterns if necessary;
 - using immutable data structures for intermediate representations;
 - high code coverage by unit and functional tests;
2. efficiency
 - keeping trade-off between execution time and memory consumption;
3. extensibility
 - clear modular architecture

5.2 Program Components

Generally, *mousquetaires* uses following ...

5.3 Parsing the input language as YTree

below: mostly mock text.

- The language-dependent syntax tree: - for now it's the C subset language which I called 'Cmin'; as a base, I used the C11 grammar from ANTLR github repository, then I simplified it a lot, cutting off many unnecessary C

syntax features and making it more convenient for parsing. When developing the Cmin language, I kept in mind C elements that are necessary for processing the linux kernel code, though for now not the whole grammar element described in file 'Cmin.g4' are being implemented; - later I am going to add the litmus grammar as well; - in future, it will be not a problem to add any new C-like language;

- The language-independent abstract syntax tree (aliased 'Ytree', where 'Y' resembles branching of the tree): - all tree nodes in my code are prefixed with 'Y', see tentative (yet almost complete) class hierarchy in picture 'YEntity.png'; - this AST contains very basic language elements according to the C execution model (statements and expressions); - converting the language-dependent syntax tree to the language-independent syntax tree is performed by Visitor pattern (e.g., for Cmin->Ytree conversion is made by 'CminToYtreeConverterVisitor') - minor changes are performed by converting to ytree representation: desugaring the target code, etc.

5.4 YTree to XGraph event converter

- Then, the AST is being interpreted and converted to event-based representation (aliased 'Xrepr' for eXecution representation): - more low-level code representation (or high-level assembly); - I try to keep this representation close to the one you described in your papers: basic load & store events, branching events, fence events; - this representation is being implementing these days, I've just started doing it (see current class hierarchy in the picture 'XEntity.png');

- After we acquired the event-based representation, we can perform some modifications/simplifications/optimisations on it (separately, allowing user to manage them): - converting to SSA form as one of necessary steps before encoding; - (more? - I'm not thinking about it yet);

5.4.1 Loop unrolling

The original program encoded into the XGraph represents a *flow graph*, a connected cyclic directed graph with single source node (ENTRY) (usually for convenience all leaves are connected to the sink node (EXIT)). The cycles are caused by low-level jump instructions, obtained from non-linear high-level control-flow statements (such as while, do-while, for, etc.). How-




Figure 5.1: Example of the flow graph from the Figure ??, unwinded up to the bound $k = 6$

ever, the cyclic flow graph cannot be encoded into SMT formula since ...
 //TODO:REFERENCE.

5.5 XGraph to ZFormula (SMT) encoder

- Then, this modified event-representation is being encoded to SMT formula and sent to the solver.

5.6 Optimisations

... performed on each stage

Chapter 6

Evaluation

6.1 Comparison with PORTHOS

6.1.1 Unique Features

6.1.2 Performance

6.2 Comparison with HERD

6.2.1 Unique Features

6.2.2 Performance

Chapter 7

Summary

Bibliography

- [LFH⁺17] H Ponce de León, Florian Furbach, Keijo Heljanko, and Roland Meyer. “Portability Analysis for Axiomatic Memory Models. PORTHOS: One Tool for all Models”. In: *CoRR* abs/1702.06704 (2017). arXiv: 1702.06704. URL: <http://arxiv.org/abs/1702.06704>.
- [McK17] Paul E McKenney. *Is parallel programming hard, and, if so, what can you do about it?(v2017. 01.02 a)*. 2017.
- [MAM⁺17] Paul E. McKenney, Jade Alglave, Luc Maranget, Andrea Parri, and Alan Stern. *A formal kernel memory-ordering model (part 1)*. 2017. URL: <https://lwn.net/Articles/718628/>.
- [WBS⁺17] John Wickerson, Mark Batty, Tyler Sorensen, and George A Constantinides. “Automatically comparing memory consistency models”. In: *ACM SIGPLAN Notices*. Vol. 52. 1. ACM. 2017, pp. 190–204.
- [ACM16] Jade Alglave, Patrick Cousot, and Luc Maranget. “Syntax and semantics of the weak consistency model specification language cat”. In: *arXiv preprint arXiv:1608.07531* (2016).
- [AMT14] Jade Alglave, Luc Maranget, and Michael Tautschnig. “Herding cats: Modelling, simulation, testing, and data mining for weak memory”. In: *ACM Transactions on Programming Languages and Systems (TOPLAS)* 36.2 (2014), p. 7.
- [Par13] Terence Parr. *The definitive ANTLR 4 reference*. Pragmatic Bookshelf, 2013.
- [CKN⁺12] Edmund M Clarke, William Klieber, Miloš Nováček, and Paolo Zuliani. “Model checking and the state explosion problem”. In: *Tools for Practical Software Verification*. Springer, 2012, pp. 1–30.

- [ISO11] ISO/IEC. “SC22/WG14. ISO/IEC 9899: 2011”. In: *Information technology – Programming languages – C*. http://www.iso.org/iso/iso_catalogue/catalogue_tc/catalogue_detail.htm (2011).
- [SSA⁺11] Susmit Sarkar, Peter Sewell, Jade Alglave, Luc Maranget, and Derek Williams. “Understanding POWER multiprocessors”. In: *ACM SIGPLAN Notices* 46.6 (2011), pp. 175–186.
- [Alg10] Jade Alglave. “A shared memory poetics”. In: *Thèse de doctorat, L’université Paris Denis Diderot* (2010).
- [AFI⁺09] Jade Alglave et al. “The semantics of Power and ARM multiprocessor machine code”. In: *Proceedings of the 4th workshop on Declarative aspects of multicore programming*. ACM. 2009, pp. 13–24.
- [MZ09] Sharad Malik and Lintao Zhang. “Boolean satisfiability from theoretical hardness to practical success”. In: *Communications of the ACM* 52.8 (2009), pp. 76–82.
- [OSS09] Scott Owens, Susmit Sarkar, and Peter Sewell. “A better x86 memory model: x86-TSO”. In: *International Conference on Theorem Proving in Higher Order Logics*. Springer. 2009, pp. 391–407.
- [MPA05] Jeremy Manson, William Pugh, and Sarita V Adve. *The Java memory model*. Vol. 40. 1. ACM, 2005.
- [CBR⁺01] Edmund Clarke, Armin Biere, Richard Raimi, and Yunshan Zhu. “Bounded model checking using satisfiability solving”. In: *Formal methods in system design* 19.1 (2001), pp. 7–34.
- [Sht00] Ofer Shtrichman. “Tuning SAT checkers for bounded model checking”. In: *International Conference on Computer Aided Verification*. Springer. 2000, pp. 480–494.
- [Lam79] Leslie Lamport. “How to make a multiprocessor computer that correctly executes multiprocess program”. In: *IEEE transactions on computers* 9 (1979), pp. 690–691.
- [Lam78] Leslie Lamport. “Time, clocks, and the ordering of events in a distributed system”. In: *Communications of the ACM* 21.7 (1978), pp. 558–565.

Appendices

Appendix A

The ANTLR grammar of input language used in PORTHOS

APPENDIX A. THE ANTLR GRAMMAR OF INPUT LANGUAGE USED IN PORTHOS 28

```

grammar Porthos;

main
:   program
;

bool_expression
:   bool_atom
|   bool_atom BOOL_OP bool_atom
;

bool_atom
:   'true'
|   'false'
|   '(' arith_expr COMP_OP arith_expr ')'
|   '(' bool_expression ')'
;

arith_expr
:   arith_atom ARITH_OP arith_atom
|   arith_atom
;

arith_atom
:   DIGIT
|   register
|   '(' arith_expr ')'
;

register
:   WORD
;

location
:   WORD
;

local
:   register '<-' arith_expr
;

load
:   register '<:-' location
;

store
:   location ':=' register
;

read
:   register '=' location '.' 'load' '('
    ATOMIC ')'
;

write
:   location '.' 'store' '(' ATOMIC ','
    register ')'
;

statement
:   expression
|   sequential_stmt
|   while_stmt
|   if_stmt
;

expression
:   local
|   load
|   store
|   FENCE
|   read
|   write
;

sequential_stmt
:   expression ';' statement
|   while_stmt ';' statement
|   if_stmt ';' statement
;

if_stmt
:   'if' bool_expression 'then' '{' statement
    '}' ('else' '{' statement '}')?
;

while_stmt
:   'while' bool_expression '{' statement '}'
;

program
:   '{' location (',' location)* '}'
    ('thread t' DIGIT '{' statement '}'
    ('exists' (location '=' DIGIT ','
    | DIGIT ':' register '=' DIGIT ','))
    )*
;

ATOMIC
:   '_na' | '_sc' | '_rx' | '_acq' | '_rel' |
    '_con'
;

FENCE
:   'mfence' | 'sync' | 'lwsync' | 'isync'
;

COMP_OP : '=' | '!=' | '<=' | '<' | '>=' |
    '>';
ARITH_OP : '+' | '-' | '*' | '/' | '%';
BOOL_OP : 'and' | 'or';

DIGIT : [0-9];
WORD : (LETTER | DIGIT)+;
LETTER : 'a'..'z' | 'A'..'Z';

```