

Aalto University
School of Science
Master's Programme in Computer, Communication and Information Sciences

Artem YUSHKOVSKIY

Automated Analysis of Weak Memory Models

Master's Thesis
Espoo, ???2018

Supervisor: Assoc. Prof. Keijo Heljanko
Instructor:

Aalto University
 School of Science

Master's Programme in Computer, Communication and In-
 formation Sciences

ABSTRACT OF
 MASTER'S THESIS

Author:	Artem YUSHKOVSKIY		
Title:	Automated Analysis of Weak Memory Models		
Date:	?.?.2018	Pages:	vi + 12
Professorship:		Code:	AS-116
Supervisor:	Assoc. Prof. Keijo Heljanko		
Instructor:			
<p>In id fringilla velit. Maecenas sed ante sit amet nisi iaculis bibendum sed vel elit. Quisque eleifend lacus nec ipsum lobortis ornare. Nam lectus diam, facilisis eget porttitor ac, fringilla quis massa. Phasellus ac dolor sem, eget varius lacus. Sed sit amet ipsum eget arcu tristique aliquam. Integer aliquam velit sit amet odio tempus commodo. Quisque commodo lacus in leo sagittis vel dignissim quam vestibulum. Cras fringilla velit et diam dictum faucibus. Pellentesque at eros non mauris auctor euismod. Nullam convallis arcu vel lectus sollicitudin rutrum. Praesent consequat, nisl at pretium posuere, neque arcu dapibus lacus, ut sollicitudin elit velit ultricies libero. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae;</p> <p>Nulla semper hendrerit molestie. Pellentesque blandit velit sit amet est vestibulum faucibus. Nullam massa turpis, venenatis non mollis fringilla, mattis et diam. Fusce molestie convallis elementum. Morbi nec lacus dapibus arcu mollis gravida. Aliquam erat volutpat. Nam vitae magna nunc. Nunc ut ipsum at massa porttitor vestibulum. Praesent diam lorem, ultrices nec vestibulum id, volutpat nec lacus.</p>			
Keywords:	Thesis template, master’s thesis		
Language:	English		

Aalto-yliopisto
 Perustieteiden korkeakoulu
 ???

???

Tekijä:	Artem YUSHKOVSKIY		
Työn nimi:	?		
Päiväys:	???.2018	Sivumäärä:	vi + 12
Professuuri:	?	Koodi:	AS-116
Valvoja:	Assoc. Prof. Keijo Heljanko		
Ohjaaja:	<p>Cras tincidunt bibendum erat, vel tincidunt diam porttitor aliquam. Donec sit amet urna non felis placerat pharetra. Aenean ultrices facilisis nulla vitae semper. Nullam non libero quis dui fermentum aliquam id vel eros. Praesent elementum tortor quis sem congue iaculis sit amet eget nisl. Quisque erat tortor, condimentum eu volutpat et, blandit et augue. Phasellus erat turpis, pretium non feugiat id, posuere id velit. Vestibulum ut sapien felis, quis convallis dui.</p> <p>In elementum est eu nulla hendrerit feugiat. In sodales diam vel lacus cursus tincidunt. Morbi nibh dui, imperdiet non vestibulum non, dignissim id risus. Sed sollicitudin neque lectus, porttitor sollicitudin elit. Nulla facilisi. Nullam in ante eu mi suscipit sollicitudin. Sed est velit, gravida facilisis varius eget, tempus sed urna. Aliquam erat volutpat. Nam semper condimentum nisi. Nullam scelerisque, metus nec sodales vulputate, purus augue venenatis urna, sit amet mattis turpis nisl ac metus. Mauris nec odio ut neque condimentum vulputate vel in turpis. Nulla facilisi. Nulla id tellus sapien, vitae blandit lorem.</p>		
Asiasanat:	Diplomityöpohja		
Kieli:	Englanti		

Acknowledgements

In id fringilla velit. Maecenas sed ante sit amet nisi iaculis bibendum sed vel elit. Quisque eleifend lacus nec ipsum lobortis ornare. Nam lectus diam, facilisis eget porttitor ac, fringilla quis massa. Phasellus ac dolor sem, eget varius lacus. Sed sit amet ipsum eget arcu tristique aliquam. Integer aliquam velit sit amet odio tempus commodo. Quisque commodo lacus in leo sagittis vel dignissim quam vestibulum. Cras fringilla velit et diam dictum faucibus. Pellentesque at eros non mauris auctor euismod. Nullam convallis arcu vel lectus sollicitudin rutrum. Praesent consequat, nisl at pretium posuere, neque arcu dapibus lacus, ut sollicitudin elit velit ultricies libero. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae;

Espoo, ???.2018

Artem YUSHKOVSKIY

Abbreviations

LI	Lorem Ipsum
ABC	Quisque et mi lacus, nec porta ante.
DEF	Proin pellentesque accumsan laoreet

Contents

Abbreviations	v
1 Introduction	1
1.1 Thesis structure	3
2 Weak Memory Models	4
2.1 The Litmus tests	5
2.2 The CAT language	5
2.3 Some known WMM	5
2.3.1 x86-TSO	5
3 Implementation	6
3.1 Program Requirements	6
3.2 Program Components	6
3.3 C11 to YTree parser	6
3.4 YTree to XGraph event converter	7
3.4.1 Loop unrolling	7
3.5 XGraph to ZFormula (SMT) encoder	9
3.6 Optimisations	9
4 Evaluation	10
4.1 Comparison with PORTHOS	10
4.1.1 Unique Features	10
4.1.2 Performance	10
4.2 Comparison with HERD	10
4.2.1 Unique Features	10
4.2.2 Performance	10
5 Summary	11

Chapter 1

Introduction

Most modern computer systems contain large parts working concurrently. Though system parallelising can improve its performance drastically, it opens numerous of problems connected to correctness, robustness and reliability, which makes the concurrent program design one of the most difficult problems of programming [1].

Traditionally, studies related to concurrent programming concern mostly on fundamental theoretical questions of designing race-free and lock-free parallel algorithms, asynchronous data structures and synchronisation primitives of a programming language. Unfortunately, when it comes to the real-world concurrent programs, the algorithmic level of abstraction is not enough for guaranteeing their properties of correctness and reliability. The reasons of this fact lie in the code optimisations that both compiler and hardware perform in order to increase performance as much as possible. For instance, Figure 1.1 provides simple example of unexpected state $(0:EAX=0 \wedge 1:EAX=0)$ reachable in x86 machines (such little examples that illustrate specific behaviours of a WMM are called *litmus tests*). This behaviour is caused by write buffers used by all processors in x86 architecture. These buffers cache writes to shared variables, so that the writes to shared memory does not become visible by other processes immediately. In the example, the write `MOV [x], 1` performed in the process P0 stores value 1 into the shared variable [x] in the write buffer of process P0. Meanwhile, the write cache of the process P1 may not have updated version of the variable [x], neither may have the main memory, so that the read `MOV EBX, [x]` performed in the process P1 may read the initial value 0, even if this variable has been already updated in another thread. More precise description of x86-TSO memory model is given in Chapter 2.3.1.

{ x=0; y=0; }	
P0	P1
MOV [x],1	MOV [y],1
MOV EAX,[y]	MOV EBX,[x]
exists (0:EAX=0 /\ 1:EAX=0)	
x86-TSO: allow	

Figure 1.1: Litmus test of memory operations reordering allowed by the x86-TSO weak memory model

The first memory model for concurrent systems was formulated by Leslie Lamport back in 1979 [2]. This memory model, called the *sequential consistency (SC) model*, allows only those executions (interleavings) that produce the same result as if the operations had been executed by single process. This means that the order of operations executed by a process is strictly defined by the program it executes. The SC model does requires the write to a shared variable performed in one process to become visible by all other processes not instantly, but simultaneously. This means each process communicates to the shared memory directly, without local buffering. Another important requirement of SC memory model is that it forbids memory operations reordering within single process (the order is strictly defined by the program).

The SC model is considered to be the strong memory model in the sence that it provides strong guarantees regarding the ordering and caused effect of memory operations. Different relaxations of this model lead to the class of *weak memory models (WMM)*. WMMs serve as set of guarantees made by designers of execution environment (hardware, programming language, compiler, database, operation system, etc.) to programmers on which behaviours of their concurrent code they may expect.

Although weak memory studies is rather young research area, there exist frameworks and tools for exploring WMMs and examining simple programs with respect to the them. The state-of-the-art tool is diy (for *do it yourself*), developed by the researchers from INRIA institute, France, and University of Cambridge, UK, and firstly released back in 2010. It is the software suite for designing and testing weak memory models, it consists of the litmus tests generators diy7, diycross7 and diyone7, the litmus tests concrete executor litmus7 that runs tests on a physical machine while collecting its behaviours, and the weak memory models simulator herd7 that implements reachability analysis for capturing states allowed by the

WMM. The latter tool uses systematic exhaustive search in a state space, which may be inefficient when analysing real-world programs, therefore more sophisticated search techniques should be applied for this problem.

One possible approach is to use an efficient implementation of an SMT-solver (a SAT-solver extended by satisfiability modulo theories) [3]. This approach allows to capture symbolically the semantics of both the program and the weak memory model while encoding it into a single SMT-formula. Most modern SMT-solvers are efficient enough to be able to operate the state space of size ??? [CITATION-?].

In the work [3], the SMT-based approach was defined for analysing the portability of a program from one hardware architecture to another, which is defined as “an execution that is consistent with the target but inconsistent with the source memory model”. Although encoding the control-flow and the data-flow of a program into an SMT-formula seems to be a trivial problem of symbolic execution, encoding of the weak memory model is more tedious. The reason is that some relations of WMMs are defined as mutually-recursive and need to be linearised in order to be encoded into an equivalent logical formula.

Current thesis aims to improve the proof-of-concept tool PORTHOS firstly introduced in April 2017 in the work [3] by extending the input language, which currently represents the minimum subset of C, and revising the general architecture of the tool in order to enhance performance, reliability and maintainability.

1.1 Thesis structure

The Chapter 2 gives more detailed description of the weak memory models analysis and provides description of memory models for some common architectures (x86, ARM and POWER, Sparc, ???). Chapter ...

Chapter 2

Weak Memory Models

Over the last decades, the problem of formalisation of the weak memory models has been developed significantly. Research of WMM aims, firstly, to formalise the weak memory models and provide systematic, sound and complete formal approach of defining WMMs in order to be able to verify systems with respect to them. //the cat language here.

Secondly, researchers work on extracting the formal hardware memory models from existing implementations or from their specifications, that are written in natural language and thus suffer from ambiguities and incompletenesses. Over last decade the memory models have been extracted for most mainstream multiprocessor architectures, such as x86-TSO (*Total Store Order*) model for x86 architecture formalised in 2009 [4], SPARC-TSO for Sparc architecture [5], much more relaxed memory model [6] for Power and ARM architectures defined in [7], Alpha [8]. Moreover, in 2005 Milner started the work on developing the weak memory model for C++ language, which was introduced in C++11 [9] standard [10] (//todo: C11 MM [6]). The memory model for Java that is based on the *happens-before* principle was introduced in JVM [11] in [12].

Thirdly, important research direction targets the problem of verifying (or at least finding bugs in) existing software systems with respect to weak memory models. In this area the notable works are on defining the Linux kernel memory model [13] (that is being actively developing these days `kernel_wmm_1`). Distributed databases also need the wmm, see transaction consistency [14].

One important question that rises naturally when talking about weak memory models is whether a program is portable from one processor architecture to another. This can be solved ...[ref to PORTHOS]

2.1 The Litmus tests

Test Execution Diagrams

2.2 The CAT language

the CAT language [ref to Jade's paper] (hard part: to decipher the Alglave's paper).

the event representation

2.3 Some known WMM

2.3.1 x86-TSO

// example with reordering // ex. with // Rev-29 Example 7-6. Stores Are Transitively Visible.

There is a barrier instruction `mfence` that may be used for flushing the buffers into the main memory.

briefly known hw memory models: X86-TSO, Alpha, POWER, – ref to Jade; language memory models: Java, C++; library-level kernel memory model, ref to github with tests

Relationship between different models http://wiki.expertiza.ncsu.edu/index.php/CSC/ECE_506_Spring_2013/10c_ks

Chapter 3

Implementation

This chapter describes the architecture of the tool *mousquitaires* ...
language: java

3.1 Program Requirements

- stability (tests)
 - scalability (new features of language, new models, new tasks for a program)
 - transparency
 - efficiency

3.2 Program Components

Big view

3.3 C11 to YTree parser

below: mostly mock text.

- The language-dependent syntax tree: - for now it's the C subset language which I called 'Cmin'; as a base, I used the C11 grammar from ANTLR github repository, then I simplified it a lot, cutting off many unnecessary C syntax features and making it more convenient for parsing. When developing the Cmin language, I kept in mind C elements that are necessary for processing the linux kernel code, though for now not the whole grammar element described in file 'Cmin.g4' are being implemented; - later I am

going to add the litmus grammar as well; - in future, it will be not a problem to add any new C-like language;

- The language-independent abstract syntax tree (aliased 'Ytree', where 'Y' resembles branching of the tree): - all tree nodes in my code are prefixed with 'Y', see tentative (yet almost complete) class hierarchy in picture 'YEntity.png'; - this AST contains very basic language elements according to the C execution model (statements and expressions); - converting the language-dependent syntax tree to the language-independent syntax tree is performed by Visitor pattern (e.g., for Cmin->Ytree conversion is made by 'CminToYtreeConverterVisitor') - minor changes are performed by converting to ytree representation: desugaring the target code, etc.

3.4 YTree to XGraph event converter

- Then, the AST is being interpreted and converted to event-based representation (aliased 'Xrepr' for eXecution representation): - more low-level code representation (or high-level assembly); - I try to keep this representation close to the one you described in your papers: basic load & store events, branching events, fence events; - this representation is being implementing these days, I've just started doing it (see current class hierarchy in the picture 'XEntity.png');

- After we acquired the event-based representation, we can perform some modifications/simplifications/optimisations on it (separately, allowing user to manage them): - converting to SSA form as one of necessary steps before encoding; - (more? - I'm not thinking about it yet);

3.4.1 Loop unrolling

The original program encoded into the XGraph represents a *flow graph*, a connected cyclic directed graph with single source node [ENTRY] (usually for convenience all leaves are connected to the sink node [EXIT]). The cycles are caused by low-level jump instructions, obtained from non-linear high-level control-flow statements (such as while, do-while, for, etc.). However, the cyclic flow graph cannot be encoded into SMT formula since ...
//TODO:REFERENCE.




Figure 3.1: Example of the flow graph from the Figure ??, unwinded up to the bound $k = 6$

3.5 XGraph to ZFormula (SMT) encoder

- Then, this modified event-representation is being encoded to SMT formula and sent to the solver.

3.6 Optimisations

... performed on each stage

Chapter 4

Evaluation

4.1 Comparison with PORTHOS

4.1.1 Unique Features

4.1.2 Performance

4.2 Comparison with HERD

4.2.1 Unique Features

4.2.2 Performance

Chapter 5

Summary

Bibliography

- [1] Paul E McKenney. *Is parallel programming hard, and, if so, what can you do about it?*(v2017. 01.02 a). 2017.
- [2] Leslie Lamport. “How to make a multiprocessor computer that correctly executes multiprocess program”. In: *IEEE transactions on computers* 9 (1979), pp. 690–691.
- [3] Hernán Ponce de León et al. “Portability Analysis for Axiomatic Memory Models. PORTHOS: One Tool for all Models”. In: *CoRR* abs/1702.06704 (2017). arXiv: 1702.06704. URL: <http://arxiv.org/abs/1702.06704>.
- [4] Scott Owens, Susmit Sarkar, and Peter Sewell. “A better x86 memory model: x86-TSO”. In: *International Conference on Theorem Proving in Higher Order Logics*. Springer. 2009, pp. 391–407.
- [5] Susmit Sarkar et al. “Understanding POWER multiprocessors”. In: *ACM SIGPLAN Notices* 46.6 (2011), pp. 175–186.
- [6] Mark Batty et al. “Mathematizing C++ concurrency”. In: *ACM SIGPLAN Notices* 46.1 (2011), pp. 55–66.
- [7] Peter Bailis et al. “Highly available transactions: Virtues and limitations”. In: *Proceedings of the VLDB Endowment* 7.3 (2013), pp. 181–192.