

Automated Analysis of Weak Memory Models

Artem Yushkovskiy^{1,2}

MSc Candidate

Supervisors: **Assoc. Prof. Keijo Heljanko¹**

Docent Igor I. Komarov²

¹**Aalto University** (Espoo, Finland)
Department of Computer Science,
School of Science

²**ITMO University** (Saint Petersburg, Russia)
Faculty of Information Security
and Computer Technologies

Espoo, Saint Petersburg, 2018

Problem statement (*Цель работы*)

To rework the proof-of-concept memory model-aware analysis tool Porthos [**Porthos17a**] by:

- ▶ extending the C-like input language,
- ▶ revising its architecture and
- ▶ re-implementing the tool in order to enhance performance, extensibility, reliability and maintainability

Task specification (Задачи работы)

- ▶ Study the general framework for memory model-aware analysis of concurrent programs [alglave2010shared];
- ▶ Review the existing tools for memory model-aware analysis;
- ▶ Examine the existing architecture of Porthos, its strengths and weaknesses;
- ▶ Design a new architecture for PorthosC that allow to extend the input language to the (large subset of) C language, be robust, transparent, efficient and extensible.

Verification of concurrent software

Motivating example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0: x \leftarrow 1$	$q_0: y \leftarrow 1$
$p_1: r_p \leftarrow y$	$q_1: r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	

Verification of concurrent software

Motivating example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0: x \leftarrow 1$	$q_0: y \leftarrow 1$
$p_1: r_p \leftarrow y$	$q_1: r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	

Single thread
(no concurrency)

p_0, p_1, q_0, q_1 (0; 1)

q_0, q_1, p_0, p_1 (1; 0)

Verification of concurrent software

Motivating example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0: x \leftarrow 1$	$q_0: y \leftarrow 1$
$p_1: r_p \leftarrow y$	$q_1: r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	

Sequential Consistency
(classical concurrency)

p_0, p_1, q_0, q_1 (0; 1)
 q_0, q_1, p_0, p_1 (1; 0)
 p_0, q_0, p_1, q_1 (1; 1)
 p_0, q_0, q_1, p_1 (1; 1)
 q_0, p_0, p_1, q_1 (1; 1)
 q_0, p_0, q_1, p_1 (1; 1)

Verification of concurrent software

Motivating example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0: x \leftarrow 1$	$q_0: y \leftarrow 1$
$p_1: r_p \leftarrow y$	$q_1: r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	

Total Store Order
(x86, Alpha, ...)

p_0, p_1, q_0, q_1 (0; 1)	$\underline{p_1}, \underline{p_0}, q_0, q_1$ (0; 1)	$p_0, p_1, \underline{q_1}, \underline{q_0}$ (0; 1)	$\underline{p_1}, \underline{p_0}, \underline{q_1}, \underline{q_0}$ (0; 1)
q_0, q_1, p_0, p_1 (1; 0)	$q_0, q_1, \underline{p_1}, \underline{p_0}$ (1; 0)	$\underline{q_1}, \underline{q_0}, p_0, p_1$ (1; 0)	$\underline{q_1}, \underline{q_0}, \underline{p_1}, \underline{p_0}$ (1; 0)
p_0, q_0, p_1, q_1 (1; 1)	$\underline{p_1}, q_0, \underline{p_0}, q_1$ (0; 1)	$p_0, \underline{q_1}, p_1, \underline{q_0}$ (0; 1)	$\underline{p_1}, \underline{q_1}, p_0, \underline{q_0}$ (0; 0)
p_0, q_0, q_1, p_1 (1; 1)	$\underline{p_1}, q_0, q_1, \underline{p_0}$ (0; 1)	$p_0, \underline{q_1}, \underline{q_0}, p_1$ (1; 1)	$\underline{p_1}, \underline{q_1}, \underline{q_0}, \underline{p_0}$ (0; 0)
q_0, p_0, p_1, q_1 (1; 1)	$q_0, \underline{p_1}, \underline{p_0}, q_1$ (1; 1)	$\underline{q_1}, p_0, p_1, \underline{q_0}$ (0; 0)	$\underline{q_1}, \underline{p_1}, \underline{p_0}, \underline{q_0}$ (0; 0)
q_0, p_0, q_1, p_1 (1; 1)	$q_0, \underline{p_1}, q_1, \underline{p_0}$ (1; 0)	$\underline{q_1}, p_0, \underline{q_0}, p_1$ (1; 0)	$\underline{q_1}, \underline{p_1}, \underline{q_0}, \underline{p_0}$ (0; 0)

Verification of concurrent software

Motivating example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0: x \leftarrow 1$	$q_0: y \leftarrow 1$
$p_1: r_p \leftarrow y$	$q_1: r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	

Total Store Order
(x86, Alpha, ...)

p_0, p_1, q_0, q_1 (0; 1)	$\underline{p_1}, \underline{p_0}, q_0, q_1$ (0; 1)	$p_0, p_1, \underline{q_1}, \underline{q_0}$ (0; 1)	$\underline{p_1}, \underline{p_0}, \underline{q_1}, \underline{q_0}$ (0; 1)
q_0, q_1, p_0, p_1 (1; 0)	$q_0, q_1, \underline{p_1}, \underline{p_0}$ (1; 0)	$\underline{q_1}, \underline{q_0}, p_0, p_1$ (1; 0)	$\underline{q_1}, \underline{q_0}, \underline{p_1}, \underline{p_0}$ (1; 0)
p_0, q_0, p_1, q_1 (1; 1)	$\underline{p_1}, q_0, \underline{p_0}, q_1$ (0; 1)	$p_0, \underline{q_1}, p_1, \underline{q_0}$ (0; 1)	$\underline{p_1}, \underline{q_1}, p_0, \underline{q_0}$ (0; 0)
p_0, q_0, q_1, p_1 (1; 1)	$\underline{p_1}, q_0, q_1, \underline{p_0}$ (0; 1)	$p_0, \underline{q_1}, \underline{q_0}, p_1$ (1; 1)	$\underline{p_1}, \underline{q_1}, \underline{q_0}, \underline{p_0}$ (0; 0)
q_0, p_0, p_1, q_1 (1; 1)	$q_0, \underline{p_1}, \underline{p_0}, q_1$ (1; 1)	$\underline{q_1}, p_0, p_1, \underline{q_0}$ (0; 0)	$\underline{q_1}, \underline{p_1}, \underline{p_0}, \underline{q_0}$ (0; 0)
q_0, p_0, q_1, p_1 (1; 1)	$q_0, \underline{p_1}, q_1, \underline{p_0}$ (1; 0)	$\underline{q_1}, p_0, \underline{q_0}, p_1$ (1; 0)	$\underline{q_1}, \underline{p_1}, \underline{q_0}, \underline{p_0}$ (0; 0)

Verification of concurrent software

Motivating example: Store buffering (hardware relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	




Figure: An x86-TSO abstract machine [sewell2010x86]

Verification of concurrent software

Motivating example: Store buffering (hardware relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	

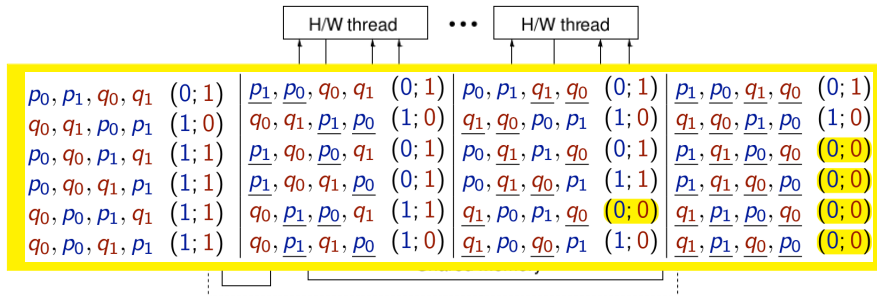


Figure: An x86-TSO abstract machine [sewell2010x86]

Weak memory model-aware analysis

Event-based program representation

- ▶ **Event** $\in \mathbb{E}$, a low-level primitive operation:
 - ▶ *memory event* $\in \mathbb{M} = \mathbb{R} \cup \mathbb{W}$: access to a local/shared memory,
 - ▶ *computational event* $\in \mathbb{C}$: computation over local memory, and
 - ▶ *barrier event* $\in \mathbb{B}$: synchronisation fences;
- ▶ **Relation** $\subseteq \mathbb{E} \times \mathbb{E}$:
 - ▶ *basic relations*:
 - ▶ *program-order* relation $\text{po} \subseteq \mathbb{E} \times \mathbb{E}$: (control-flow),
 - ▶ *read-from* relation $\text{rf} \subseteq \mathbb{W} \times \mathbb{R}$: (data-flow), and
 - ▶ *coherence-order* relation $\text{co} \subseteq \mathbb{W} \times \mathbb{W}$: (data-flow);
 - ▶ *derived relations*:
 - ▶ *union* $\text{r1} \mid \text{r2}$,
 - ▶ *sequence* $\text{r1} ; \text{r2}$,
 - ▶ *transitive closure* r^+ ,
 - ▶ \dots ;
- ▶ **Assertion** over relations or sets of events:
 - ▶ *acyclicity*, *irreflexivity* or *emptiness*

Weak memory model-aware analysis

Testing candidate executions




Figure: The four candidate executions allowed under x86-TSO

Weak memory model-aware analysis

Testing candidate executions

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	

SC model:

...

$\text{fr} = (\text{rf}^{-1}; \text{co})$

$\text{acyclic}(\text{fr} \cup \text{po})$




Figure: The four candidate executions allowed under x86-TSO

Portability analysis

The Porthos tool

```
/* Prints Hello World */  
#include <stdio.h>  
int main (void) {  
    printf ("Hello World!");  
    return 0;  
}
```




Figure: The illustration of the portability problem [Porthos17slides]

Portability analysis

The Porthos tool

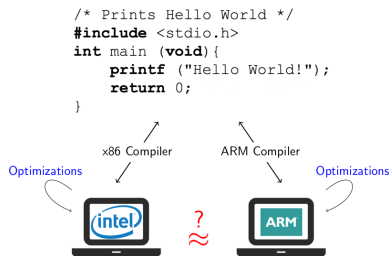


Figure: The illustration of the portability problem [Porthos17slides]

Definition (Portability [Porthos17a])

The program P is portable from the source memory model \mathcal{M}_S to the target one \mathcal{M}_T if

$$\text{cons}_{\mathcal{M}_T}(P) \subseteq \text{cons}_{\mathcal{M}_S}(P)$$

- Portability as an SMT-based bounded reachability problem:
 $\phi = \phi_{CF} \wedge \phi_{DF} \wedge \phi_{\mathcal{M}_T} \wedge \phi_{\neg \mathcal{M}_S}$
- $\text{SAT}(\phi) \implies \text{portability bug}$

Portability analysis

The Porthos tool




Figure: The illustration of the portability problem [Porthos17slides]

Definition

(Portability [Porthos17a])

The program P is portable from the source memory model \mathcal{M}_S to the target one \mathcal{M}_T if

$$\text{cons}_{\mathcal{M}_T}(P) \subseteq \text{cons}_{\mathcal{M}_S}(P)$$

- Portability as an SMT-based bounded reachability problem:
 $\phi = \phi_{CF} \wedge \phi_{DF} \wedge \phi_{\mathcal{M}_T} \wedge \phi_{\neg \mathcal{M}_S}$
- $\text{SAT}(\phi) \implies$ portability bug

Portability analysis

Porthos v1: An input example

```
{ flag0, flag1, turn }

thread t0 {
  while true {
    a <- 1;
    b <- 0;
    flag0.store(_rx, a);
    f1 = flag1.load(_rx);
    while (f1 == 1) {
      t1 = turn.load(_rx);
      if (t1 != 0) {
        flag0.store(_rx,b);
        t1 = turn.load(_rx);
        while (t1 != 0) {
          t1 = turn.load(_rx)
        };
        flag0.store(_rx,a)
      }
    }
  }
}

thread t1 {
  ...
}

exists turn=10,
```

Portability analysis

Porthos v1: An input example


```
{ flag0, flag1, turn }

thread t0 {
  while true {
    a <- 1;
    b <- 0;
    flag0.store(_rx, a);
    f1 = flag1.load(_rx);
    while (f1 == 1) {
      t1 = turn.load(_rx);
      if (t1 != 0) {
        flag0.store(_rx,b);
        t1 = turn.load(_rx);
        while (t1 != 0) {
          t1 = turn.load(_rx)
        };
        flag0.store(_rx,a)
      }
    }
  }
}

thread t1 {
  ...
}

exists turn=10,
```

Encoding for the control-flow: An example



$$\begin{aligned}
 \phi_{CF} = & [\mathbf{x}(e_2) \Rightarrow \mathbf{x}(e_1)] \\
 & \wedge [\mathbf{x}(e_3) \Rightarrow \mathbf{x}(e_2)] \\
 & \wedge [\mathbf{x}(e_{nop_1}) \Rightarrow \mathbf{x}(e_1)] \\
 & \wedge [\mathbf{x}(e_{nop_2}) \Rightarrow \mathbf{x}(e_2)] \\
 & \wedge [\mathbf{x}(e_4) \Rightarrow (\mathbf{x}(e_{nop_1}) \vee \mathbf{x}(e_3) \vee \mathbf{x}(e_{nop_2}))] \\
 & \wedge [\mathbf{x}(e_{nop_1}) \wedge \mathbf{x}(e_1) \Rightarrow g_{1,4}] \\
 & \wedge [(\mathbf{x}(e_3) \wedge \mathbf{x}(e_2)) \Rightarrow g_{2,3}] \\
 & \wedge [(\mathbf{x}(e_{nop_2}) \wedge \mathbf{x}(e_2)) \Rightarrow g_{2,4}] \\
 & \wedge \neg[\mathbf{x}(e_2) \wedge \mathbf{x}(e_{nop_1})] \\
 & \wedge \neg[\mathbf{x}(e_3) \wedge \mathbf{x}(e_{nop_2})]
 \end{aligned}$$

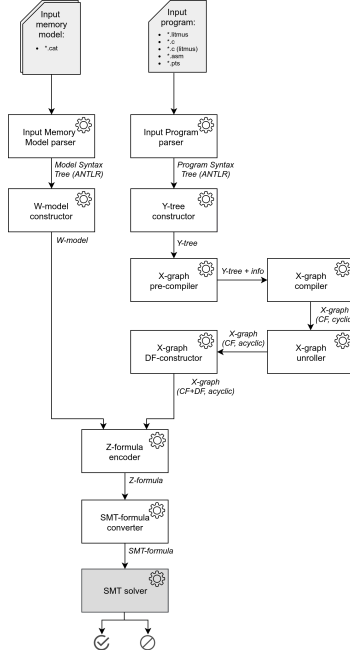
Figure: Example of encoding for the control-flow of the event-flow graph

The input language

The input language parser used by Porthos suffered from several disadvantages:

- ▶ it contained the parser code inlined directly into the grammar (hardly maintainable);
- ▶ the semantics of operations and kinds of variables (global or shared) were determined syntactically (4 different types of assignment: '=', ':=', '<- ' and '<:- ', each for different kinds of arguments);
- ▶ restricted syntax for expressions.
- ▶ In contrast, PorthosC uses the full C language grammar of proposed in the C11 standard [jtc2011sc22] and the visitor that converts the ANTLR grammar to the AST (Y-tree).

Architecture



The X-graph internal representation




Figure: The inheritance tree of main X-graph interfaces

The X-graph compiler

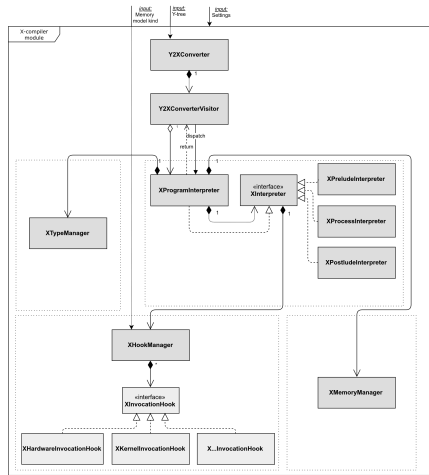


Figure: Main components of the X-compilation processing unit

X-graph unrolling




Figure: Example of the flow graph unrolling up to bound $k = 6$

Evaluation

Much better.
[to be done]

Summary

- ▶ The general framework for memory model-aware analysis was implemented in PorthosC;
- ▶ The input language has been extended;
- ▶ The old architecture of Porthos has been analysed and considered while designing the new architecture for PorthosC;
- ▶ to be done: more

Bibliography I