

# Automated Analysis of Weak Memory Models

**Artem Yushkovskiy<sup>1,2</sup>**

MSc Candidate

Supervisors: **Assoc. Prof. Keijo Heljanko<sup>1</sup>**

**Docent Igor I. Komarov<sup>2</sup>**

<sup>1</sup>Department of Computer Science,  
School of Science,  
**Aalto University** (Espoo, Finland)

<sup>2</sup>Faculty of Information Security  
and Computer Technologies,  
**ITMO University** (Saint Petersburg, Russia)

# Outline

## Introduction

The problem

The background: weak memory model-aware analysis

Portability Analysis

SMT-encoding

## Implementation

The input language and general architecture

The X-graph: structure and construction

## Evaluation



# Outline

## Introduction

The problem

The background: weak memory model-aware analysis

Portability Analysis

SMT-encoding

## Implementation

The input language and general architecture

The X-graph: structure and construction

## Evaluation

## Problem statement (Цель работы)

To rework the proof-of-concept memory model-aware analysis tool Porthos [3] by:

- extending the C-like input language,
- revising its architecture and
- re-implementing the tool in order to enhance performance, extensibility, reliability and maintainability



## Task specification (Задачи работы)

- Study the general framework for memory model-aware analysis of concurrent programs [1];
- Review existing tools for memory model-aware analysis;
- Investigate existing architecture of Porthos, its strengths and weaknesses;
- Design a new architecture for PorthosC that *allow to* easily support the C input language, be robust, transparent, efficient and extensible.



# Verification of concurrent software

Example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$



# Verification of concurrent software

Example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$

SC

$p_0, p_1, q_0, q_1$  (0; 1)  
 $q_0, q_1, p_0, p_1$  (1; 0)



# Verification of concurrent software

Example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$

SC

$p_0, p_1, q_0, q_1$  (0; 1)  
 $q_0, q_1, p_0, p_1$  (1; 0)  
 $p_0, q_0, p_1, q_1$  (1; 1)  
 $p_0, q_0, q_1, p_1$  (1; 1)  
 $q_0, p_0, p_1, q_1$  (1; 1)  
 $q_0, p_0, q_1, p_1$  (1; 1)



# Verification of concurrent software

Example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$

SC

$p_0, p_1, q_0, q_1$	(0; 1)
$q_0, q_1, p_0, p_1$	(1; 0)
$p_0, q_0, p_1, q_1$	(1; 1)
$p_0, q_0, q_1, p_1$	(1; 1)
$q_0, p_0, p_1, q_1$	(1; 1)
$q_0, p_0, q_1, p_1$	(1; 1)

TSO

$\underline{p_1}, \underline{p_0}, q_0, q_1$	(0; 1)
$q_0, q_1, \underline{p_1}, \underline{p_0}$	(1; 0)
$\underline{p_1}, q_0, \underline{p_0}, q_1$	(0; 1)
$\underline{p_1}, q_0, q_1, \underline{p_0}$	(0; 0)
$q_0, \underline{p_1}, \underline{p_0}, q_1$	(1; 1)
$q_0, \underline{p_1}, q_1, \underline{p_0}$	(1; 0)
$p_0, p_1, \underline{q_1}, \underline{q_0}$	(0; 1)
$\underline{q_1}, \underline{q_0}, p_0, p_1$	(1; 0)
$p_0, \underline{q_1}, p_1, \underline{q_0}$	(0; 1)
$p_0, \underline{q_1}, \underline{q_0}, p_1$	(1; 1)
$\underline{q_1}, p_0, p_1, \underline{q_0}$	(0; 0)
$\underline{q_1}, p_0, \underline{q_0}, p_1$	(1; 0)
$\underline{p_1}, \underline{p_0}, \underline{q_1}, \underline{q_0}$	(0; 1)
$\underline{q_1}, \underline{q_0}, \underline{p_1}, \underline{p_0}$	(1; 0)
$\underline{p_1}, \underline{q_1}, \underline{p_0}, \underline{q_0}$	(0; 0)
$\underline{p_1}, \underline{q_1}, \underline{q_0}, \underline{p_0}$	(0; 0)
$\underline{q_1}, \underline{p_1}, \underline{p_0}, \underline{q_0}$	(0; 0)
$\underline{q_1}, \underline{p_1}, \underline{q_0}, \underline{p_0}$	(0; 0)

# Verification of concurrent software

Example: Store buffering (hardware relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$




Figure: An x86-TSO abstract machine [4]

# The weak memory model

## Axiomatic semantics: The definition

- **Event**  $\in \mathbb{E}$ , a low-level primitive operation:
  - *memory event*  $\in \mathbb{M} = \mathbb{R} \cup \mathbb{W}$ : access to a local/shared memory,
  - *computational event*  $\in \mathbb{C}$ : computation over local memory, and
  - *barrier event*  $\in \mathbb{B}$ : synchronisation fences;
- **Relation**  $\subseteq \mathbb{E} \times \mathbb{E}$ :
  - *basic relations*:
    - *program-order* relation  $po \subseteq \mathbb{E} \times \mathbb{E}$ : (control-flow),
    - *read-from* relation  $rf \subseteq \mathbb{W} \times \mathbb{R}$ : (data-flow), and
    - *coherence-order* relation  $co \subseteq \mathbb{W} \times \mathbb{W}$ : (data-flow);
  - *derived relations*:
    - *union*  $r1 \mid r2$ ,
    - *sequence*  $r1 ; r2$ ,
    - *transitive closure*  $r^+$ ,
    - $\dots$ ;
- **Assertion** over relations or sets of events:
  - *acyclicity*, *irreflexivity* or *emptiness*



# The weak memory model

## Testing the candidate executions

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$




Figure: The four candidate executions allowed under x86-TSO

# The weak memory model

## Testing the candidate executions

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$

SC model:

...

$fr = (rf^{-1}; co)$

$acyclic(fr \cup po)$




Figure: The four candidate executions allowed under x86-TSO

## Tools for memory model-aware analysis

- diy tool suite:
  - diy, diycross and diyone, litmus tests generators,
  - litmus, a litmus test concrete executor, and
  - herd, a weak memory model simulator;
- the stateless model checkers (CHES, Nidhugg);
- the tool for automated synthesis of the synchronisation primitives musketeer;
- the instrumenting compiler goto-cc which is a part of CBMC model checker;
- the tool Porthos for analysing the portability of the C programs;
- and others.

# Portability analysis

## The Porthos tool

- Let the function  $cons_{\mathcal{M}}(P)$  calculate the set of executions of program  $P$  consistent under the memory model  $\mathcal{M}$ .


### Definition (Portability [3])

Let  $\mathcal{M}_{\mathcal{S}}$ ,  $\mathcal{M}_{\mathcal{T}}$  be two weak memory models. The program  $P$  is portable from  $\mathcal{M}_{\mathcal{S}}$  to  $\mathcal{M}_{\mathcal{T}}$  if  $cons_{\mathcal{M}_{\mathcal{T}}}(P) \subseteq cons_{\mathcal{M}_{\mathcal{S}}}(P)$

- Portability as an SMT-based bounded reachability problem:  
 $\phi = \phi_{CF} \wedge \phi_{DF} \wedge \phi_{\mathcal{M}_{\mathcal{T}}} \wedge \phi_{\neg \mathcal{M}_{\mathcal{S}}}$
- $SAT(\phi) \implies$  the portability bug

## Encoding for the control-flow

- Porthos v1 used another encoding scheme, where the high-level instructions were represented in the SMT-formula by separate variables:  $\phi_{CF}(i_2; i_3) = (cf_{i_1} \Leftrightarrow (cf_{i_2} \wedge cf_{i_3})) \wedge \phi_{CF}(i_2) \wedge \phi_{CF}(i_3)$ .
- In PorthosC, the high-level AST firstly is compiled into the *event-flow graph* with events as nodes and relations as edges.
- All edges are labelled by *guards*, local-memory computations ( $\epsilon$  denotes an empty guard).



(a) The sequence      (b) Conditional branching      (c) Branch merging

**Figure:** Possible mutual arrangements of events in a control-flow graph



## Encoding for the control-flow

- Let  $\mathbf{x} : \mathbb{E} \rightarrow \{0, 1\}$  be the predicate that signifies the fact that the event has been executed.
- The control-flow of the program is encoded as following:

$$\phi_{CF_{seq}} = \mathbf{x}(e_{i+1}) \Rightarrow \mathbf{x}(e_i)$$

$$\begin{aligned} \phi_{CF_{br}} = & [\mathbf{x}(e_i) \Rightarrow \mathbf{x}(e_k)] \wedge \cdots \wedge [\mathbf{x}(e_{i+j}) \Rightarrow \mathbf{x}(e_k)] \\ & \wedge [\mathbf{x}(e_i) \wedge \mathbf{x}(e_k) \Rightarrow g_{i,k}] \wedge \cdots \wedge [\mathbf{x}(e_{i+j}) \wedge \mathbf{x}(e_k) \Rightarrow g_{i+j,k}] \\ & \wedge \cdots \end{aligned}$$

$$\wedge \left( \bigvee_{e_l \in \text{succ}(e_m)} \bigvee_{\substack{e_n \in \text{succ}(e_k) \\ e_n \neq e_m}} \neg[\mathbf{x}(e_m) \wedge \mathbf{x}(e_n)] \right)$$

$$\phi_{CF_{mer}} = \mathbf{x}(e_k) \Rightarrow \left( \bigvee_{e_p \in \text{pred}(e_k)} \mathbf{x}(e_p) \right)$$

## Encoding for the data-flow

- SSA-indices are computed as following:
  - any access to a shared variable (both read and write) increments its SSA-index;
  - only writes to a local variable increment its SSA-index (reads preserve indices);
  - no access to a constant variable or computed (evaluated) expression changes their SSA-index.

The data-flow of an event is encoded as following:

$$\phi_{DF_{e=\text{load}(r \leftarrow l)}} = [\mathbf{x}(e) \Rightarrow (r_{i+1} = l_{i+1})]$$

$$\phi_{DF_{e=\text{store}(l \leftarrow r)}} = [\mathbf{x}(e) \Rightarrow (l_{i+1} = r_i)]$$

$$\phi_{DF_{e=\text{eval}(\cdot)}} = [\mathbf{x}(e) \Rightarrow \mathbf{v}(e)]$$

$$\phi_{DF_{mem}}(e_1, e_2) = [\text{rf}(e_1, e_2) \Rightarrow (l_i = l_j)]$$

# Outline

## Introduction

The problem

The background: weak memory model-aware analysis

Portability Analysis

SMT-encoding

## Implementation

The input language and general architecture

The X-graph: structure and construction

## Evaluation



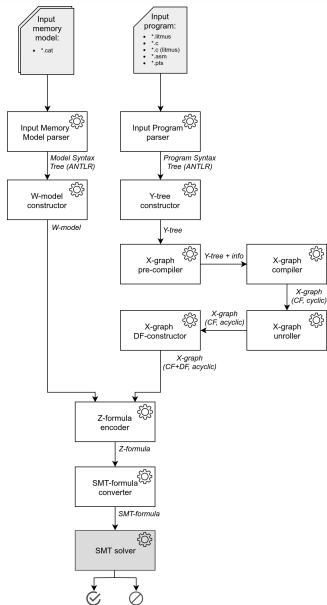
## The input language

The input language parser used by Porthos suffered from several disadvantages:

- it contained the parser code inlined directly into the grammar (hardly maintainable);
- the semantics of operations and kinds of variables (global or shared) were determined syntactically (4 different types of assignment: '=', ':=', '<-' and '<:-', each for different kinds of arguments);
- restricted syntax for expressions.
- In contrast, PorthosC uses the full C language grammar of proposed in the C11 standard [2] and the visitor that converts the ANTLR grammar to the AST (Y-tree).



# Architecture



# The X-graph internal representation


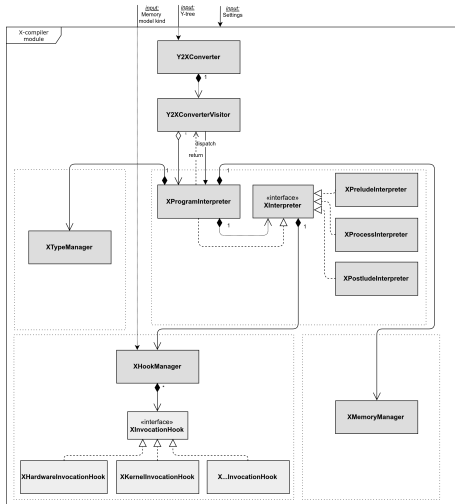



Figure: The inheritance tree of main X-graph interfaces

## The X-graph compiler



**Figure:** Main components of the X-compilation processing unit

# X-graph unrolling



**Figure:** Example of the flow graph unrolling up to bound  $k = 6$



# Outline

## Introduction

The problem

The background: weak memory model-aware analysis

Portability Analysis

SMT-encoding

## Implementation

The input language and general architecture

The X-graph: structure and construction

## Evaluation



# Evaluation




[to be done]



## Summary

- The general framework for memory model-aware analysis was implemented in PorthosC;
- The input language has been extended;
- The old architecture of Porthos has been analysed and considered while designing the new architecture for PorthosC;
- to be done: more

# Bibliography I

-  Jade Alglave. “A shared memory poetics”. In: *La Thèse de doctorat, L'université Paris Denis Diderot* (2010).
-  ISO/IEC. SC22/WG14. *ISO/IEC 9899: 2011*. Tech. rep. Geneva, Switzerland, 2011.
-  Hernán Ponce de León et al. “Portability Analysis for Weak Memory Models. PORTHOS: One Tool for all Models”. In: *Static Analysis - 24th International Symposium, SAS 2017, New York, NY, USA, August 30 - September 1, 2017, Proceedings*. 2017, pp. 299–320. DOI: 10.1007/978-3-319-66706-5\_15. URL: [https://doi.org/10.1007/978-3-319-66706-5\\_15](https://doi.org/10.1007/978-3-319-66706-5_15).

## Bibliography II



Peter Sewell et al. “x86-TSO: a rigorous and usable programmer’s model for x86 multiprocessors”. In: *Communications of the ACM* 53.7 (2010), pp. 89–97.