

Automated Analysis of Weak Memory Models

Artem Yushkovskiy^{1,2}

MSc Candidate

Supervisors: **Assoc. Prof. Keijo Heljanko¹**

Docent Igor I. Komarov²

¹Department of Computer Science,
School of Science,
Aalto University (Espoo, Finland)

²Faculty of Information Security
and Computer Technologies,
ITMO University (Saint Petersburg, Russia)

Espoo, Saint Petersburg, 2018

Problem statement (Цель работы)

To rework the proof-of-concept memory model-aware analysis tool Porthos [2] by:

- extending the C-like input language,
- revising its architecture and
- re-implementing the tool in order to enhance performance, extensibility, reliability and maintainability

Task specification (Задачи работы)

- Study the general framework for memory model-aware analysis of concurrent programs [1];
- Review the existing tools for memory model-aware analysis;
- Examine the existing architecture of Porthos, its strengths and weaknesses;
- Design a new architecture for PorthosC that allow to extend the input language to the (large subset of) C language, be robust, transparent, efficient and extensible.

Verification of concurrent software

Example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists $(r_p = 0 \wedge r_q = 0)$	

Verification of concurrent software

Example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists $(r_p = 0 \wedge r_q = 0)$	

Sequential
Consistency

p_0, p_1, q_0, q_1 (0; 1)
 q_0, q_1, p_0, p_1 (1; 0)

Verification of concurrent software

Example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists $(r_p = 0 \wedge r_q = 0)$	

Sequential
Consistency

p_0, p_1, q_0, q_1 (0; 1)
 q_0, q_1, p_0, p_1 (1; 0)
 p_0, q_0, p_1, q_1 (1; 1)
 p_0, q_0, q_1, p_1 (1; 1)
 q_0, p_0, p_1, q_1 (1; 1)
 q_0, p_0, q_1, p_1 (1; 1)

Verification of concurrent software

Example: Write-write reordering (compiler relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists $(r_p = 0 \wedge r_q = 0)$	

Sequential
Consistency

Total Store Order

p_0, p_1, q_0, q_1 (0; 1)	$\underline{p_1}, \underline{p_0}, q_0, q_1$ (0; 1)	$p_0, p_1, \underline{q_1}, \underline{q_0}$ (0; 1)	$\underline{p_1}, \underline{p_0}, \underline{q_1}, \underline{q_0}$ (0; 1)
q_0, q_1, p_0, p_1 (1; 0)	$q_0, q_1, \underline{p_1}, \underline{p_0}$ (1; 0)	$\underline{q_1}, \underline{q_0}, p_0, p_1$ (1; 0)	$\underline{q_1}, \underline{q_0}, \underline{p_1}, \underline{p_0}$ (1; 0)
p_0, q_0, p_1, q_1 (1; 1)	$\underline{p_1}, q_0, \underline{p_0}, q_1$ (0; 1)	$p_0, \underline{q_1}, p_1, \underline{q_0}$ (0; 1)	$\underline{p_1}, \underline{q_1}, \underline{p_0}, \underline{q_0}$ (0; 0)
p_0, q_0, q_1, p_1 (1; 1)	$\underline{p_1}, q_0, q_1, \underline{p_0}$ (0; 1)	$p_0, \underline{q_1}, \underline{q_0}, p_1$ (1; 1)	$\underline{p_1}, \underline{q_1}, \underline{q_0}, \underline{p_0}$ (0; 0)
q_0, p_0, p_1, q_1 (1; 1)	$q_0, \underline{p_1}, \underline{p_0}, q_1$ (1; 1)	$\underline{q_1}, p_0, p_1, \underline{q_0}$ (0; 0)	$\underline{q_1}, \underline{p_1}, \underline{p_0}, \underline{q_0}$ (0; 0)
q_0, p_0, q_1, p_1 (1; 1)	$q_0, \underline{p_1}, q_1, \underline{p_0}$ (1; 0)	$\underline{q_1}, p_0, \underline{q_0}, p_1$ (1; 0)	$\underline{q_1}, \underline{p_1}, \underline{q_0}, \underline{p_0}$ (0; 0)

Verification of concurrent software

Example: Store buffering (hardware relaxations)

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists $(r_p = 0 \wedge r_q = 0)$	




Figure: An x86-TSO abstract machine [3]

The weak memory model

Axiomatic semantics: The definition

- **Event** $\in \mathbb{E}$, a low-level primitive operation:
 - *memory event* $\in \mathbb{M} = \mathbb{R} \cup \mathbb{W}$: access to a local/shared memory,
 - *computational event* $\in \mathbb{C}$: computation over local memory, and
 - *barrier event* $\in \mathbb{B}$: synchronisation fences;
- **Relation** $\subseteq \mathbb{E} \times \mathbb{E}$:
 - *basic relations*:
 - *program-order* relation $\text{po} \subseteq \mathbb{E} \times \mathbb{E}$: (control-flow),
 - *read-from* relation $\text{rf} \subseteq \mathbb{W} \times \mathbb{R}$: (data-flow), and
 - *coherence-order* relation $\text{co} \subseteq \mathbb{W} \times \mathbb{W}$: (data-flow);
 - *derived relations*:
 - *union* $\text{r1} \mid \text{r2}$,
 - *sequence* $\text{r1} ; \text{r2}$,
 - *transitive closure* r^+ ,
 - \dots ;
- **Assertion** over relations or sets of events:
 - *acyclicity*, *irreflexivity* or *emptiness*

The weak memory model

Testing candidate executions

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	




Figure: The four candidate executions allowed under x86-TSO

The weak memory model

Testing candidate executions

{ x=0; y=0; }	
P	Q
$p_0 : x \leftarrow 1$	$q_0 : y \leftarrow 1$
$p_1 : r_p \leftarrow y$	$q_1 : r_q \leftarrow x$
exists ($r_p = 0 \wedge r_q = 0$)	

SC model:

...
 $fr = (rf^{-1}; co)$
 $acyclic(fr \cup po)$




Figure: The four candidate executions allowed under x86-TSO

Tools for memory model-aware analysis

- diy tool suite:
 - diy, diycross and diyone, litmus tests generators,
 - litmus, a litmus test concrete executor, and
 - herd, a weak memory model simulator;
- the stateless model checkers (CHESS, Nidhugg);
- the tool for automated synthesis of the synchronisation primitives musketeer;
- the instrumenting compiler goto-cc which is a part of CBMC model checker;
- the tool Porthos for analysing the portability of the C programs;
- and others.

Portability analysis


The Porthos tool

- Let the function $cons_{\mathcal{M}}(P)$ calculate the set of executions of program P consistent under the memory model \mathcal{M} .

Definition (Portability [2])

Let \mathcal{M}_S , \mathcal{M}_T be two weak memory models. The program P is portable from \mathcal{M}_S to \mathcal{M}_T if $cons_{\mathcal{M}_T}(P) \subseteq cons_{\mathcal{M}_S}(P)$

- Portability as an SMT-based bounded reachability problem:
 $\phi = \phi_{CF} \wedge \phi_{DF} \wedge \phi_{\mathcal{M}_T} \wedge \phi_{\neg \mathcal{M}_S}$
- $SAT(\phi) \implies$ the portability bug



$$\begin{aligned} \phi_{CF} = & [\mathbf{x}(e_2) \Rightarrow \mathbf{x}(e_1)] \\ & \wedge [\mathbf{x}(e_3) \Rightarrow \mathbf{x}(e_2)] \\ & \wedge [\mathbf{x}(e_{nop_1}) \Rightarrow \mathbf{x}(e_1)] \\ & \wedge [\mathbf{x}(e_{nop_2}) \Rightarrow \mathbf{x}(e_2)] \\ & \wedge [\mathbf{x}(e_4) \Rightarrow (\mathbf{x}(e_{nop_1}) \vee \mathbf{x}(e_3) \vee \mathbf{x}(e_{nop_2}))] \\ & \wedge [\mathbf{x}(e_{nop_1}) \wedge \mathbf{x}(e_1) \Rightarrow g_{1,4}] \\ & \wedge [(\mathbf{x}(e_3) \wedge \mathbf{x}(e_2)) \Rightarrow g_{2,3}] \\ & \wedge [(\mathbf{x}(e_{nop_2}) \wedge \mathbf{x}(e_2)) \Rightarrow g_{2,4}] \\ & \wedge \neg[\mathbf{x}(e_2) \wedge \mathbf{x}(e_{nop_1})] \\ & \wedge \neg[\mathbf{x}(e_3) \wedge \mathbf{x}(e_{nop_2})] \end{aligned}$$

Encoding for the data-flow

- SSA-indices are computed as following:
 - any access to a shared variable (both read and write) increments its SSA-index;
 - only writes to a local variable increment its SSA-index (reads preserve indices);
 - no access to a constant variable or computed (evaluated) expression changes their SSA-index.

The data-flow of an event is encoded as following:

$$\begin{aligned}
 \phi_{DF_{e=\text{load}(r \leftarrow l)}} &= [\mathbf{x}(e) \Rightarrow (r_{i+1} = l_{i+1})] \\
 \phi_{DF_{e=\text{store}(l \leftarrow r)}} &= [\mathbf{x}(e) \Rightarrow (l_{i+1} = r_i)] \\
 \phi_{DF_{e=\text{eval}(\cdot)}} &= [\mathbf{x}(e) \Rightarrow \mathbf{v}(e)]
 \end{aligned}
 \tag{1}$$

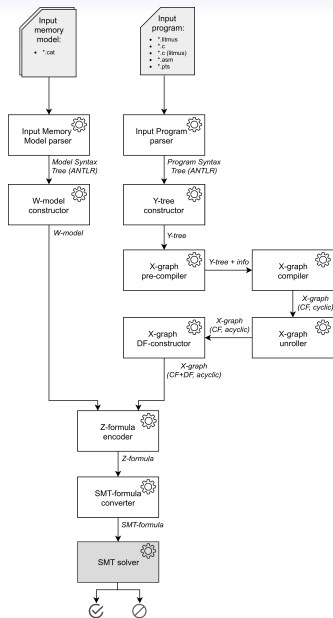
The input language

The input language parser used by Porthos suffered from several disadvantages:

- it contained the parser code inlined directly into the grammar (hardly maintainable);
- the semantics of operations and kinds of variables (global or shared) were determined syntactically (4 different types of assignment: '=', ':=', '<-' and '<:-', each for different kinds of arguments);
- restricted syntax for expressions.
- In contrast, PorthosC uses the full C language grammar of proposed in the C11 standard [jtc2011sc22] and the visitor that converts the ANTLR grammar to the AST (Y-tree).



Architecture



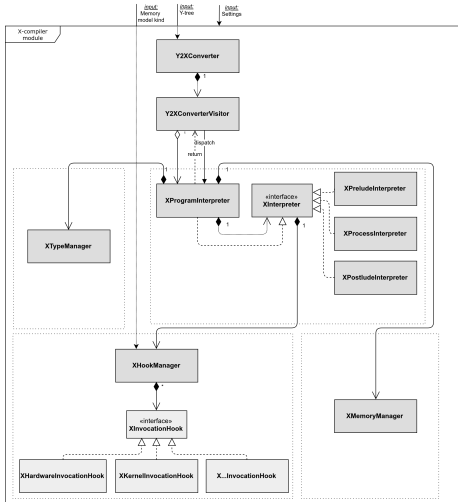



Figure: Main components of the X-compilation processing unit



Evaluation

Much better.

[to be done]

Summary

- The general framework for memory model-aware analysis was implemented in PorthosC;
- The input language has been extended;
- The old architecture of Porthos has been analysed and considered while designing the new architecture for PorthosC;
- to be done: more

