

Automated Analysis of Weak Memory Models

Artem Yushkovskiy^{1,2}

MSc Candidate

Supervisors: **Assoc. Prof. Keijo Heljanko¹**

Docent Igor I. Komarov²

¹Department of Computer Science,
School of Science,
Aalto University (Espoo, Finland)

²Faculty of Information Security
and Computer Technologies,
ITMO University (Saint Petersburg, Russia)

Espoo, Saint Petersburg, 2018

Outline

Outline

Problem statement (Цель работы)

To rework the proof-of-concept memory model-aware analysis tool Porthos [**Porthos17a**] by:

- extending the C-like input language,
- revising its architecture and
- re-implementing the tool in order to enhance performance, extensibility, reliability and maintainability

Task specification (Задачи работы)

- Study the general framework for memory model-aware analysis of concurrent programs [alglave2010shared];
- Review existing tools for memory model-aware analysis;
- Investigate existing architecture of Porthos, its strengths and weaknesses;
- Design a new architecture for PorthosC that *allow to* easily support the C input language, be robust, transparent, efficient and extensible.

Outline

Verification of Concurrent Software

Example: Write-write reordering (compiler relaxations)

| { x=0; y=0; } | |
|--------------------------|--------------------------|
| P | Q |
| $p_0 : x \leftarrow 1$ | $q_0 : y \leftarrow 1$ |
| $p_1 : r_p \leftarrow y$ | $q_1 : r_q \leftarrow x$ |

Verification of Concurrent Software

Example: Write-write reordering (compiler relaxations)

| { x=0; y=0; } | |
|--------------------------|--------------------------|
| P | Q |
| $p_0 : x \leftarrow 1$ | $q_0 : y \leftarrow 1$ |
| $p_1 : r_p \leftarrow y$ | $q_1 : r_q \leftarrow x$ |

SC

$p_0, p_1, q_0, q_1 \ (0; 1)$
 $q_0, q_1, p_0, p_1 \ (1; 0)$

Verification of Concurrent Software

Example: Write-write reordering (compiler relaxations)

| { x=0; y=0; } | |
|--------------------------|--------------------------|
| P | Q |
| $p_0 : x \leftarrow 1$ | $q_0 : y \leftarrow 1$ |
| $p_1 : r_p \leftarrow y$ | $q_1 : r_q \leftarrow x$ |

SC

p_0, p_1, q_0, q_1 (0; 1)
 q_0, q_1, p_0, p_1 (1; 0)
 p_0, q_0, p_1, q_1 (1; 1)
 p_0, q_0, q_1, p_1 (1; 1)
 q_0, p_0, p_1, q_1 (1; 1)
 q_0, p_0, q_1, p_1 (1; 1)

Verification of Concurrent Software

Example: Write-write reordering (compiler relaxations)

| { x=0; y=0; } | |
|--------------------------|--------------------------|
| P | Q |
| $p_0 : x \leftarrow 1$ | $q_0 : y \leftarrow 1$ |
| $p_1 : r_p \leftarrow y$ | $q_1 : r_q \leftarrow x$ |

SC

| | |
|----------------------|--------|
| p_0, p_1, q_0, q_1 | (0; 1) |
| q_0, q_1, p_0, p_1 | (1; 0) |
| p_0, q_0, p_1, q_1 | (1; 1) |
| p_0, q_0, q_1, p_1 | (1; 1) |
| q_0, p_0, p_1, q_1 | (1; 1) |
| q_0, p_0, q_1, p_1 | (1; 1) |

TSO

| | | | | | |
|--|--------|--|--------|--|--------|
| $\underline{p_1}, \underline{p_0}, q_0, q_1$ | (0; 1) | $p_0, p_1, \underline{q_1}, \underline{q_0}$ | (0; 1) | $\underline{p_1}, \underline{p_0}, \underline{q_1}, \underline{q_0}$ | (0; 1) |
| $q_0, q_1, \underline{p_1}, \underline{p_0}$ | (1; 0) | $\underline{q_1}, \underline{q_0}, p_0, p_1$ | (1; 0) | $\underline{q_1}, \underline{q_0}, \underline{p_1}, \underline{p_0}$ | (1; 0) |
| $\underline{p_1}, q_0, \underline{p_0}, q_1$ | (0; 1) | $p_0, \underline{q_1}, p_1, \underline{q_0}$ | (0; 1) | $\underline{p_1}, \underline{q_1}, \underline{p_0}, \underline{q_0}$ | (0; 0) |
| $\underline{p_1}, q_0, q_1, \underline{p_0}$ | (0; 0) | $p_0, \underline{q_1}, \underline{q_0}, p_1$ | (1; 1) | $\underline{p_1}, \underline{q_1}, \underline{q_0}, \underline{p_0}$ | (0; 0) |
| $q_0, \underline{p_1}, \underline{p_0}, q_1$ | (1; 1) | $\underline{q_1}, p_0, p_1, \underline{q_0}$ | (0; 0) | $\underline{q_1}, \underline{p_1}, \underline{p_0}, \underline{q_0}$ | (0; 0) |
| $q_0, \underline{p_1}, q_1, \underline{p_0}$ | (1; 0) | $\underline{q_1}, p_0, \underline{q_0}, p_1$ | (1; 0) | $\underline{q_1}, \underline{p_1}, \underline{q_0}, \underline{p_0}$ | (0; 0) |

Verification of Concurrent Software

Example: Store buffering (by hardware)

| { x=0; y=0; } | |
|--------------------------|--------------------------|
| P | Q |
| $p_0 : x \leftarrow 1$ | $q_0 : y \leftarrow 1$ |
| $p_1 : r_p \leftarrow y$ | $q_1 : r_q \leftarrow x$ |




Figure: An x86-TSO abstract machine [sewell2010x86]

Weak Memory Model

Axiomatic semantics

- **Event** $\in \mathbb{E}$, a low-level primitive operation:
 - *memory event* $\in \mathbb{M} = \mathbb{R} \cup \mathbb{W}$: access to a local/shared memory,
 - *computational event* $\in \mathbb{C}$: computation over local memory, and
 - *barrier event* $\in \mathbb{B}$: synchronisation fences;
- **Relation** $\subseteq \mathbb{E} \times \mathbb{E}$:
 - *basic relations*:
 - *program-order* relation $\text{po} \subseteq \mathbb{E} \times \mathbb{E}$: (control-flow),
 - *read-from* relation $\text{rf} \subseteq \mathbb{W} \times \mathbb{R}$: (data-flow), and
 - *coherence-order* relation $\text{co} \subseteq \mathbb{W} \times \mathbb{W}$: (data-flow);
 - *derived relations*:
 - *union* $\text{r1} \mid \text{r2}$,
 - *sequence* $\text{r1} ; \text{r2}$,
 - *transitive closure* r^+ ,
 - \dots ;
- **Assertion** over relations or sets of events:
 - *acyclicity, irreflexivity or emptiness*

Weak Memory Model

Example: Store buffering (hardware relaxations)

| { x=0; y=0; } | |
|--------------------------|--------------------------|
| P | Q |
| $p_0 : x \leftarrow 1$ | $q_0 : y \leftarrow 1$ |
| $p_1 : r_p \leftarrow y$ | $q_1 : r_q \leftarrow x$ |

SC model:

...

$\text{fr} = (\text{rf}^{-1}; \text{co})$

$\text{acyclic}(\text{fr} \cup \text{po})$




Figure: The four candidate executions allowed under x86-TSO

Outline

Portability analysis

The Porthos tool

- Let the function $cons_{\mathcal{M}}(P)$ calculate the set of executions of program P consistent under the memory model \mathcal{M} .

Definition (Portability [Porthos17a])

Let $\mathcal{M}_{\mathcal{S}}$, $\mathcal{M}_{\mathcal{T}}$ be two weak memory models. The program P is portable from $\mathcal{M}_{\mathcal{S}}$ to $\mathcal{M}_{\mathcal{T}}$ if $cons_{\mathcal{M}_{\mathcal{T}}}(P) \subseteq cons_{\mathcal{M}_{\mathcal{S}}}(P)$

- Portability as an SMT-based bounded reachability problem:
$$\phi = \phi_{CF} \wedge \phi_{DF} \wedge \phi_{\mathcal{M}_{\mathcal{T}}} \wedge \phi_{\neg \mathcal{M}_{\mathcal{S}}}$$
- $SAT(\phi) \implies$ the portability bug


Outline

Encoding for the control-flow

- In Porthos, the high-level instructions were represented in the SMT-formula by separate variables.
- For example, the sequential instruction $i_1 = i_2; i_3$ was encoded as $\phi_{CF}(i_2; i_3) = (cf_{i_1} \Leftrightarrow (cf_{i_2} \wedge cf_{i_3})) \wedge \phi_{CF}(i_2) \wedge \phi_{CF}(i_3)$,

Encoding for the control-flow

- The new encoding scheme of PorthosC follows [heljanko2008unfoldings] in general.
- The high-level AST is compiled into the *event-flow graph*, where edges (transitions) are labelled by the *guard* (ϵ is an empty guard).



(a) The sequence (b) Conditional branching (c) Branch merging

Figure: Possible mutual arrangements of events in a control-flow graph

Encoding for the control-flow

The control-flow is encoded as following:

$$\phi_{CF_{seq}} = \mathbf{x}(e_{i+1}) \Rightarrow \mathbf{x}(e_i)$$

$$\begin{aligned} \phi_{CF_{br}} = & [\mathbf{x}(e_i) \Rightarrow \mathbf{x}(e_k)] \wedge \cdots \wedge [\mathbf{x}(e_{i+j}) \Rightarrow \mathbf{x}(e_k)] \\ & \wedge [\mathbf{x}(e_i) \wedge \mathbf{x}(e_k) \Rightarrow g_{i,k}] \wedge \cdots \wedge [\mathbf{x}(e_{i+j}) \wedge \mathbf{x}(e_k) \Rightarrow g_{i+j,k}] \\ & \wedge \cdots \\ & \wedge \left(\bigvee_{e_l \in \text{succ}(e_m)} \bigvee_{\substack{e_n \in \text{succ}(e_k) \\ e_n \neq e_m}} \neg[\mathbf{x}(e_m) \wedge \mathbf{x}(e_n)] \right) \end{aligned}$$

$$\phi_{CF_{mer}} = \mathbf{x}(e_k) \Rightarrow \left(\bigvee_{e_p \in \text{pred}(e_k)} \mathbf{x}(e_p) \right)$$

Encoding for the data-flow

- SSA-indices are computed as following:
 - any access to a shared variable (both read and write) increments its SSA-index;
 - only writes to a local variable increment its SSA-index (reads preserve indices);
 - no access to a constant variable or computed (evaluated) expression changes their SSA-index.

The data-flow of an event is encoded as following:

$$\phi_{DF_{e=\text{load}(r \leftarrow l)}} = [\mathbf{x}(e) \Rightarrow (r_{i+1} = l_{i+1})]$$

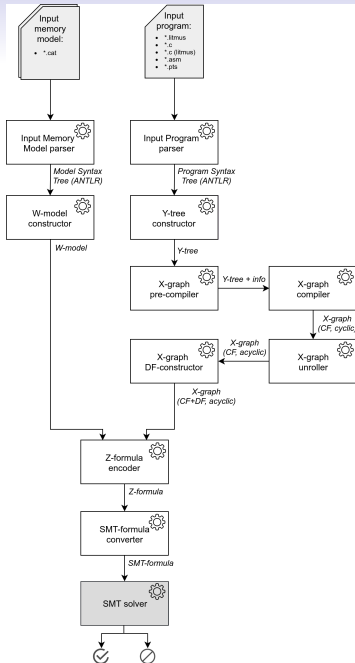
$$\phi_{DF_{e=\text{store}(l \leftarrow r)}} = [\mathbf{x}(e) \Rightarrow (l_{i+1} = r_i)]$$

$$\phi_{DF_{e=\text{eval}(\cdot)}} = [\mathbf{x}(e) \Rightarrow \mathbf{v}(e)]$$

$$\phi_{DF_{mem}}(e_1, e_2) = [\text{rf}(e_1, e_2) \Rightarrow (l_i = l_j)]$$

Outline

Architecture



The input language

The input language parser used by Porthos suffered from several disadvantages:

- it contained the parser code inlined directly into the grammar (hardly maintainable);
- the semantics of operations and kinds of variables (global or shared) were determined syntactically (4 different types of assignment: '=', ':=', '<-' and '<:-', each for different kinds of arguments);
- restricted syntax for expressions.
- In contrast, PorthosC uses the full C language grammar of proposed in the C11 standard [jtc2011sc22] and the visitor that converts the ANTLR grammar to the AST (Y-tree).

The X-graph




Figure: The inheritance tree of main X-graph interfaces

The X-graph compiler

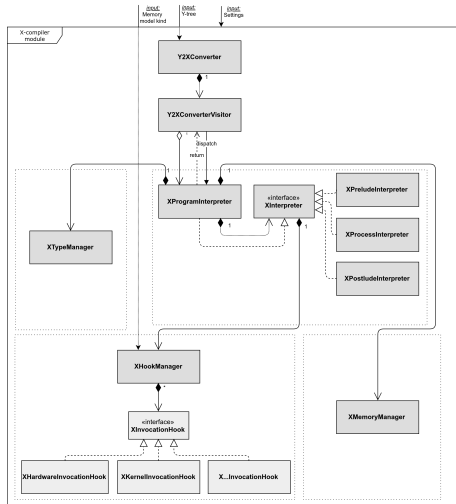


Figure: Main components of the X-compilation processing unit

X-graph unrolling




Figure: Example of the flow graph unrolling up to bound $k = 6$

Summary

- The **first main message** of your talk in one or two lines.
- The **second main message** of your talk in one or two lines.
- Perhaps a **third message**, but not more than that.
- Outlook
 - Something you haven't solved.
 - Something else you haven't solved.

Bibliography I